

Graph Algorithms II

Summer 2018 • Lecture 07/24

A Few Notes

Midterm Review Session office hours

Today 12-1:20 p.m. in STLC 115

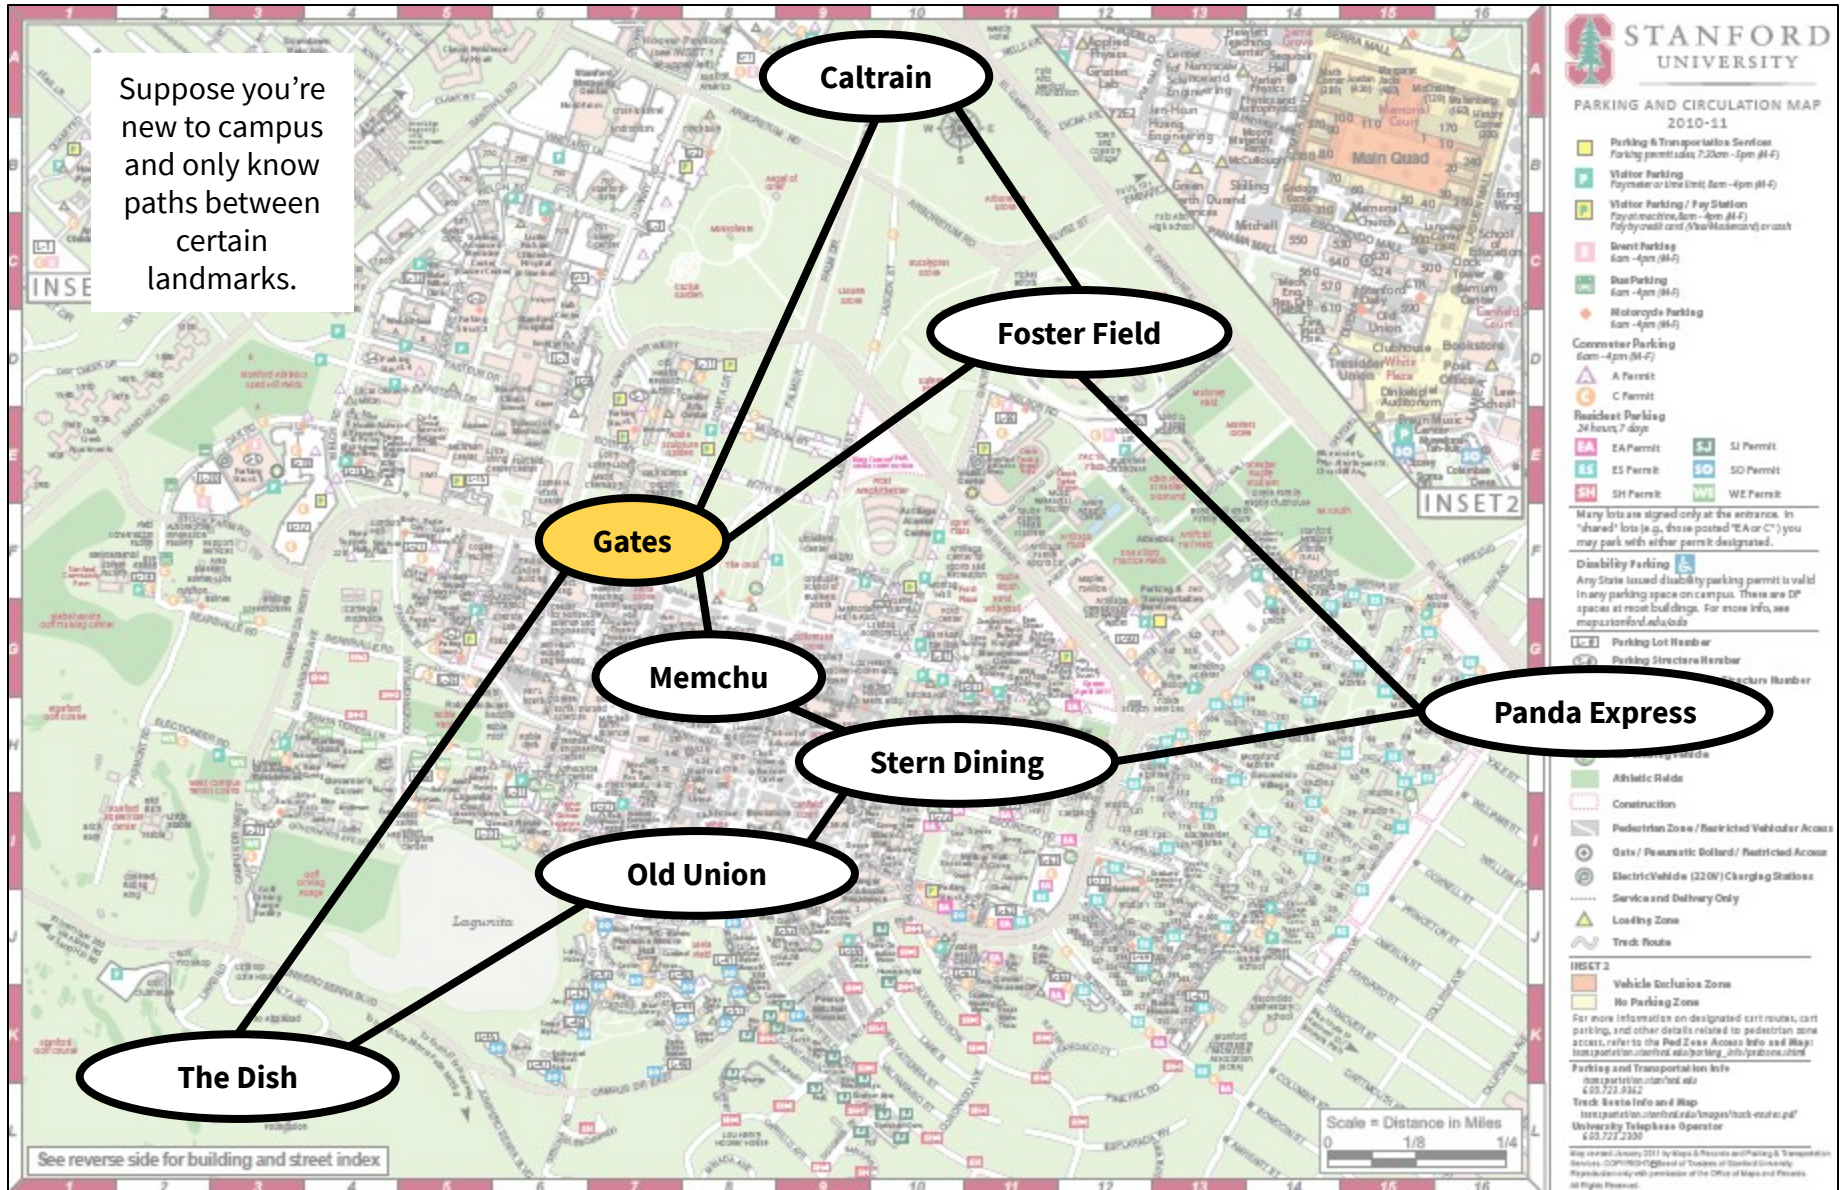
Homework 3

Due today (**you can't use late days!**)

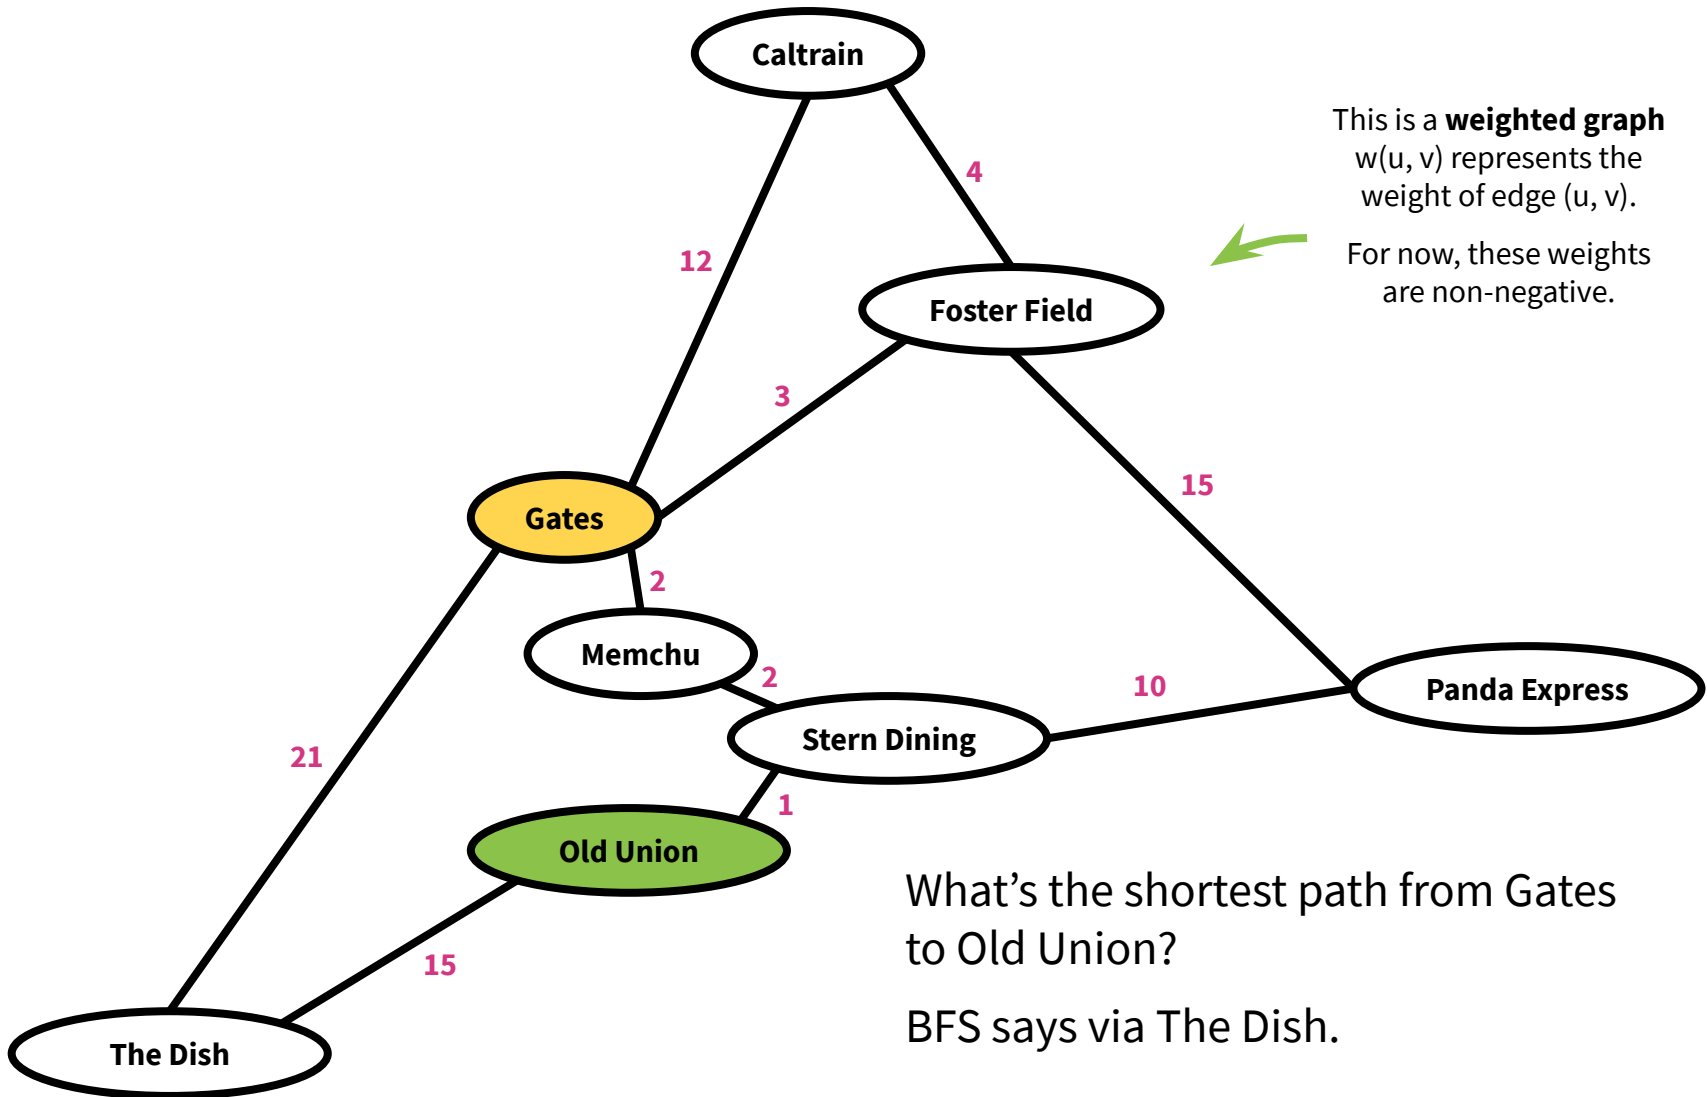
Dijkstra's Algorithm

Shortest Path

Suppose you're new to campus and only know paths between certain landmarks.



Shortest Path

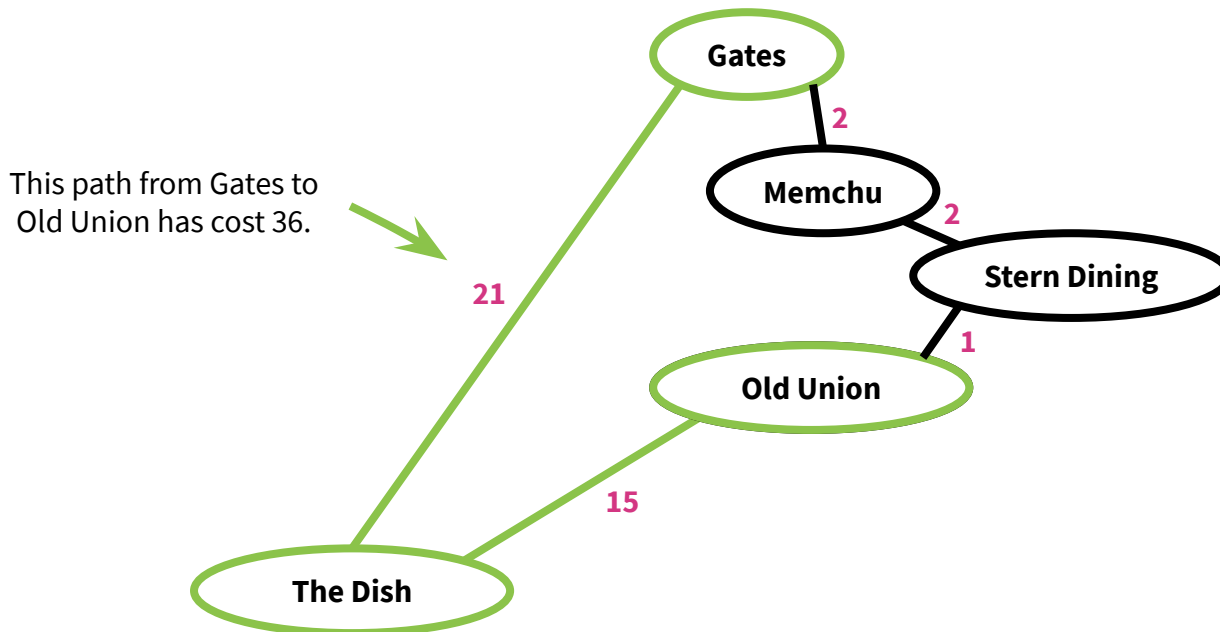


Shortest Path

What is the **shortest path** between u and v in a weighted graph?

The cost of a path is the sum of the weights along that path.

The shortest path is the one with the minimum cost.

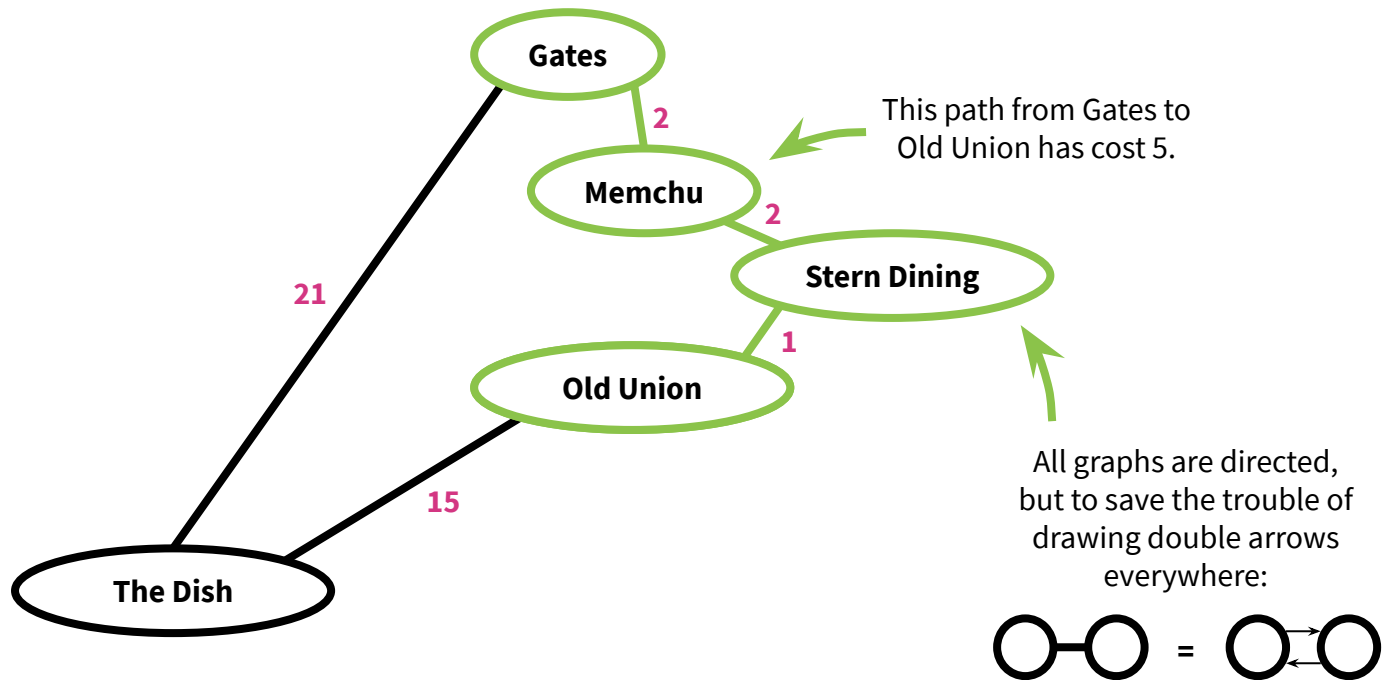


Shortest Path

What is the **shortest path** between u and v in a weighted graph?

The cost of a path is the sum of the weights along that path.


The shortest path is the one with the minimum cost.



Shortest Path

Claim: A subpath of a shortest path is also a shortest path.

Intuition:




Suppose **this** is a shortest path from **s** to **t**.

Shortest Path

Claim: A subpath of a shortest path is also a shortest path.

Intuition:



Suppose **this** is a shortest path from **s** to **t**.


Then **this** is a shortest path from **s** to **x**.

Why? 🤔

Shortest Path

Claim: A subpath of a shortest path is also a shortest path.

Intuition:



Suppose **this** is a shortest path from **s** to **t**.


Then **this** is a shortest path from **s** to **x**.

Why? 🤔

Shortest Path

Claim: A subpath of a shortest path is also a shortest path.

Intuition:



Suppose **this** is a shortest path from **s** to **t**.


Then **this** is a shortest path from **s** to **x**.

Why? 🤔 By contradiction, suppose there exists a shorter path from **s** to **x**, namely **this** one.

Shortest Path

Claim: A subpath of a shortest path is also a shortest path.

Intuition:



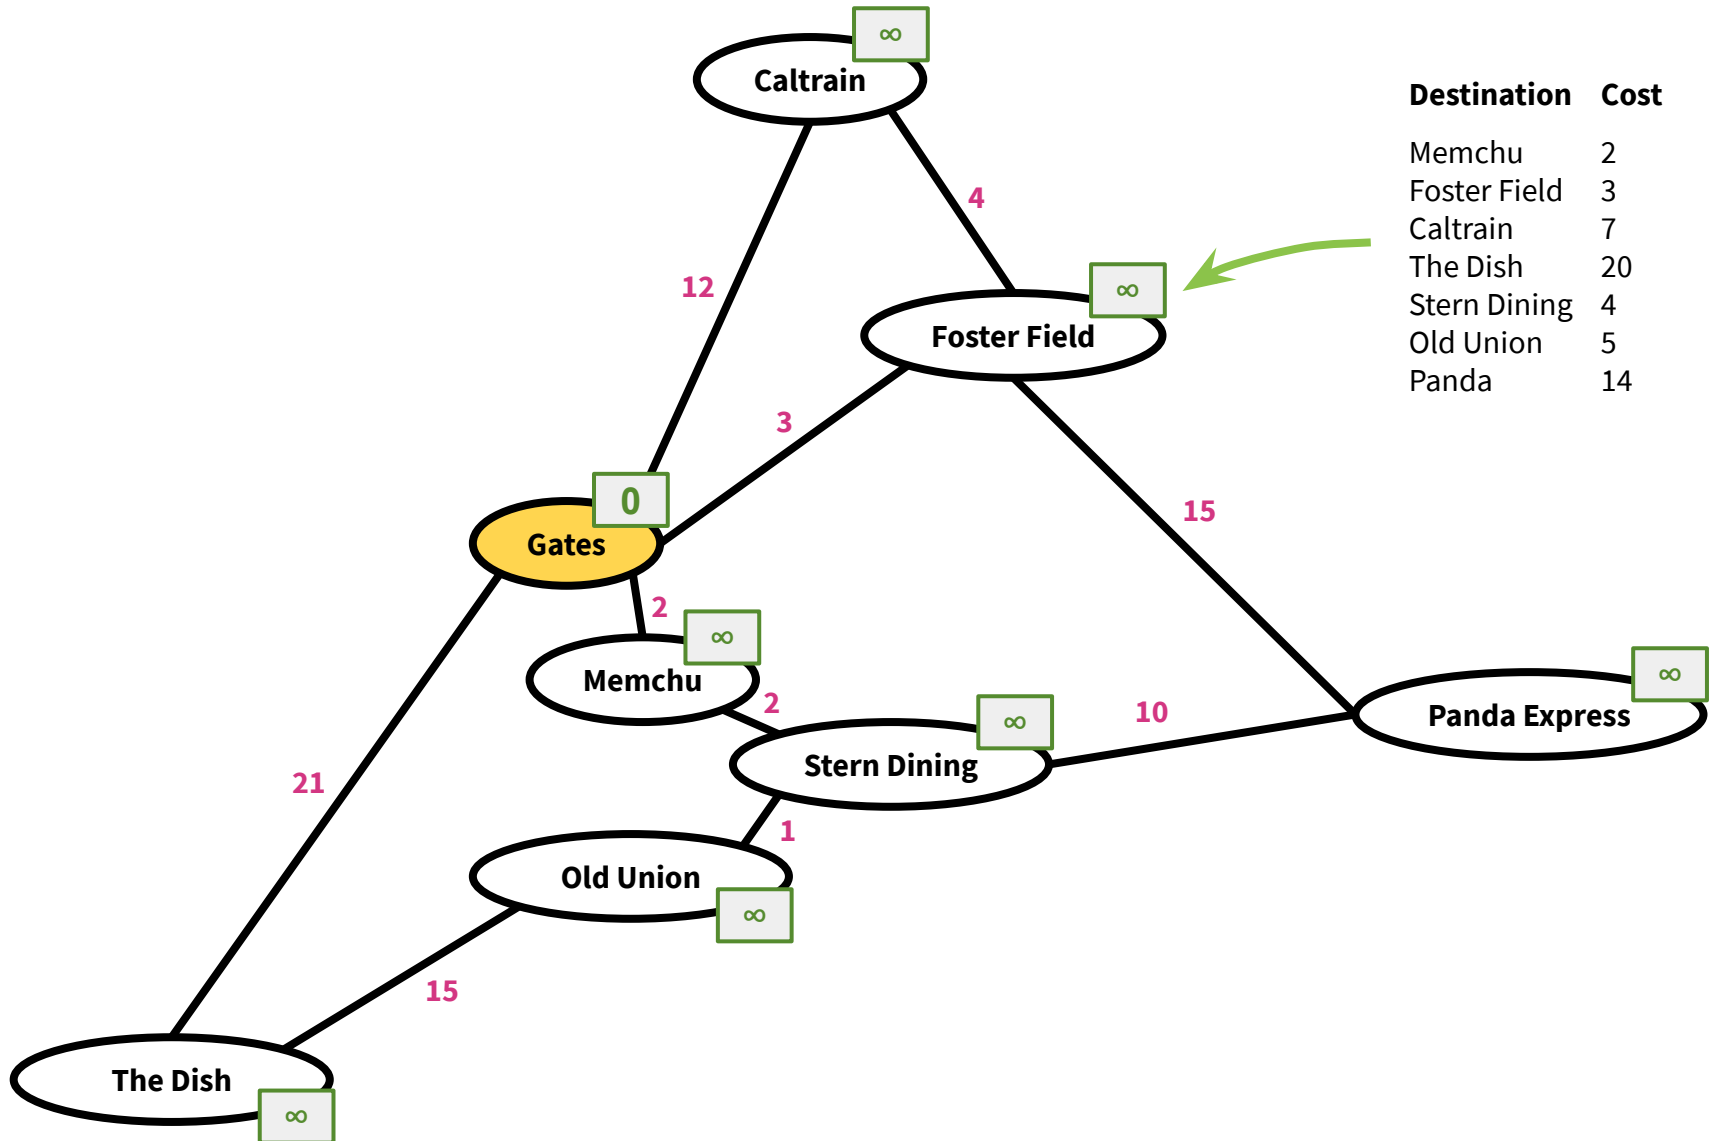
Suppose **this** is a shortest path from **s** to **t**.

Then **this** is a shortest path from **s** to **x**.

Why? 🤔 By contradiction, suppose there exists a shorter path from **s** to **x**, namely **this** one.

But then **this** is shorter than **this** shortest path from **s** to **t**.

Single-Source Shortest Path



Single-Source Shortest Path

Application: Finding the shortest path from Palo Alto to [somewhere else] for a commuter using BART, Caltrain, bike, walking, Uber, Lyft, etc.

Edge weights are a function of time, money, hassle that change depending on the commuter's mood on that day.

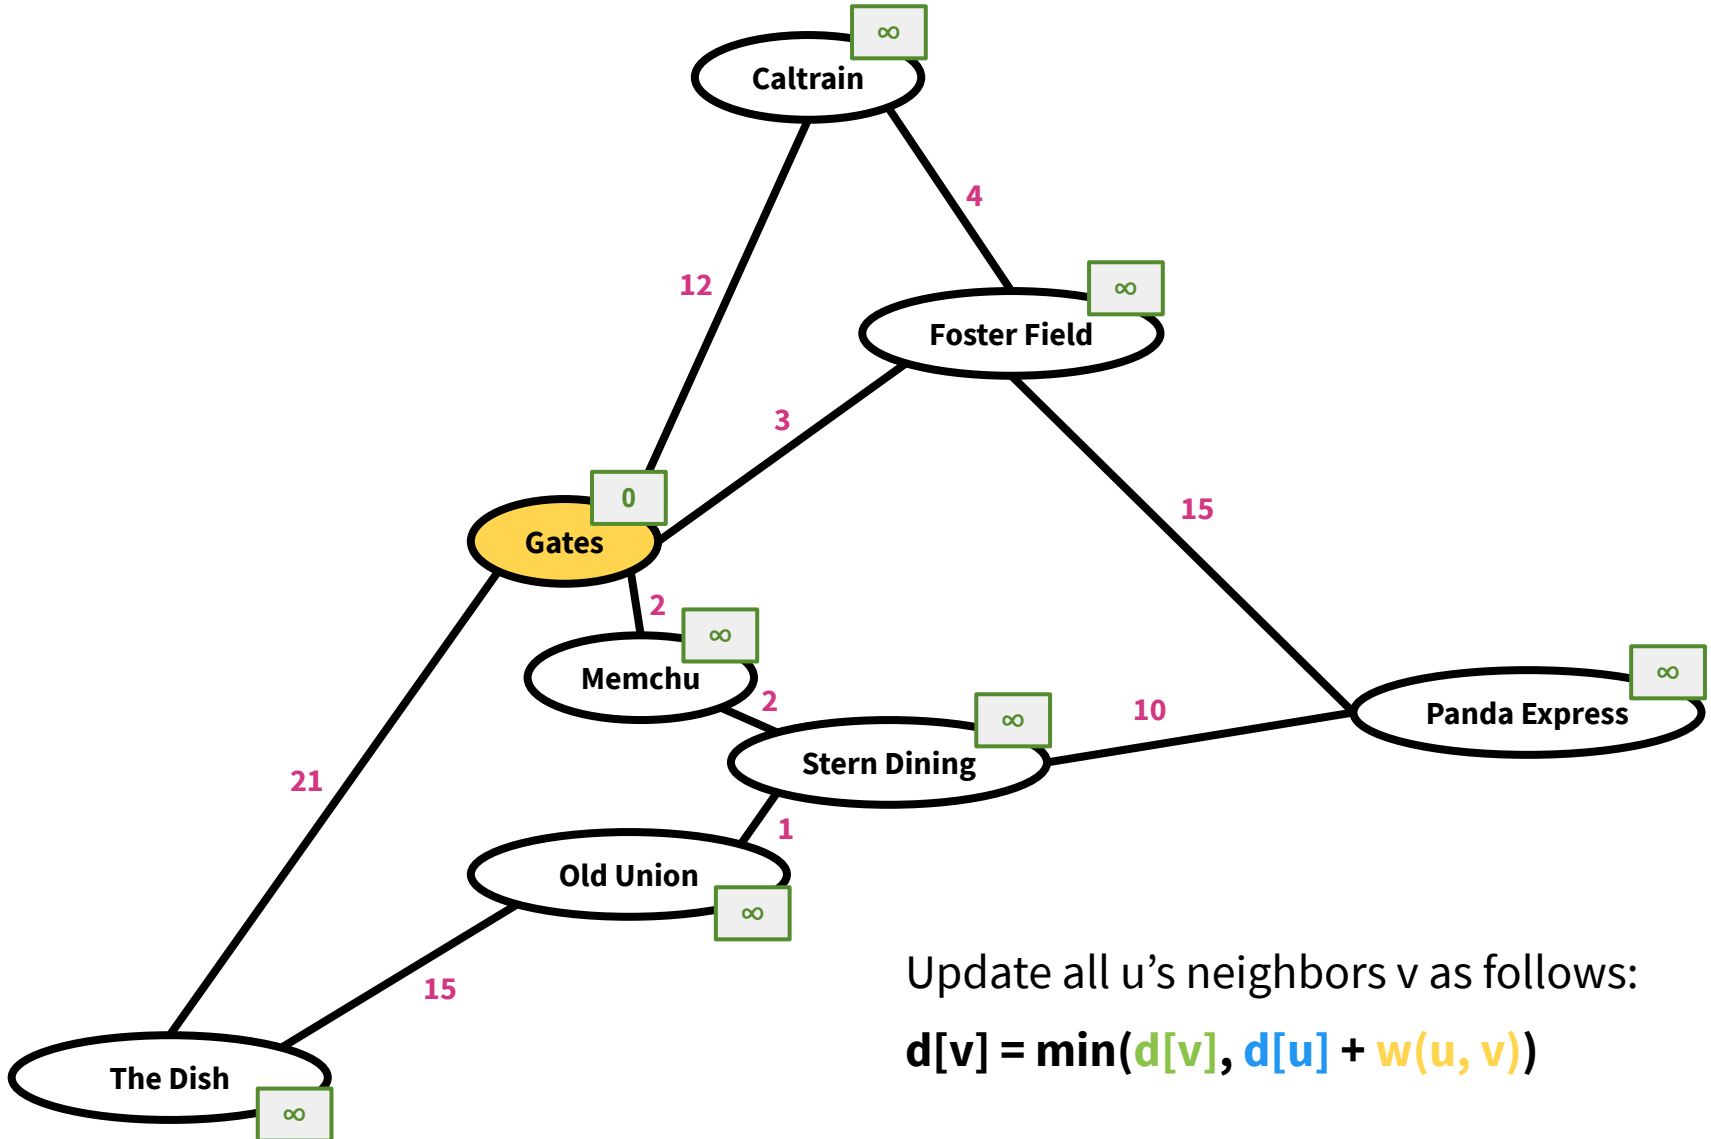
Application: Finding the shortest path from my computer to the desired server for packets using the Internet.

Edge weights are a function of link length, traffic, other costs, etc.

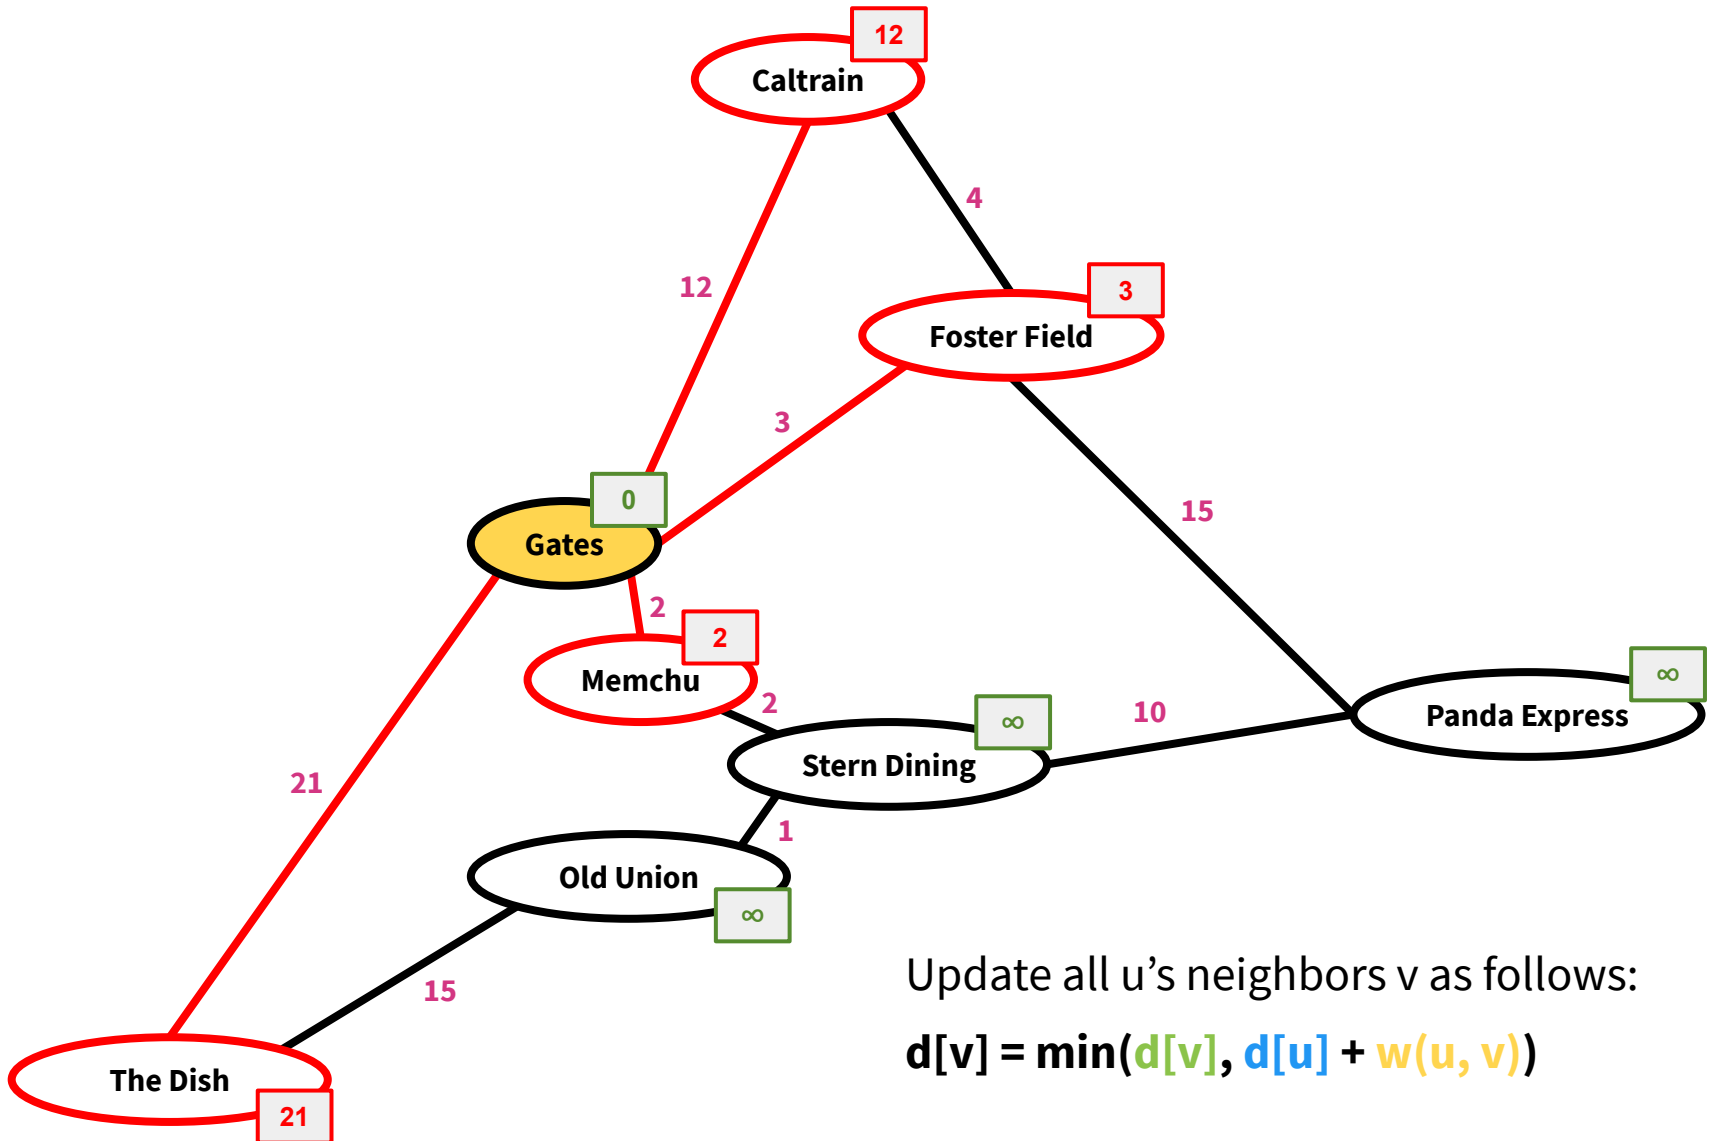
Dijkstra's Algorithm

Dijkstra's Algorithm solves the single-source shortest path problem.

Dijkstra's Algorithm



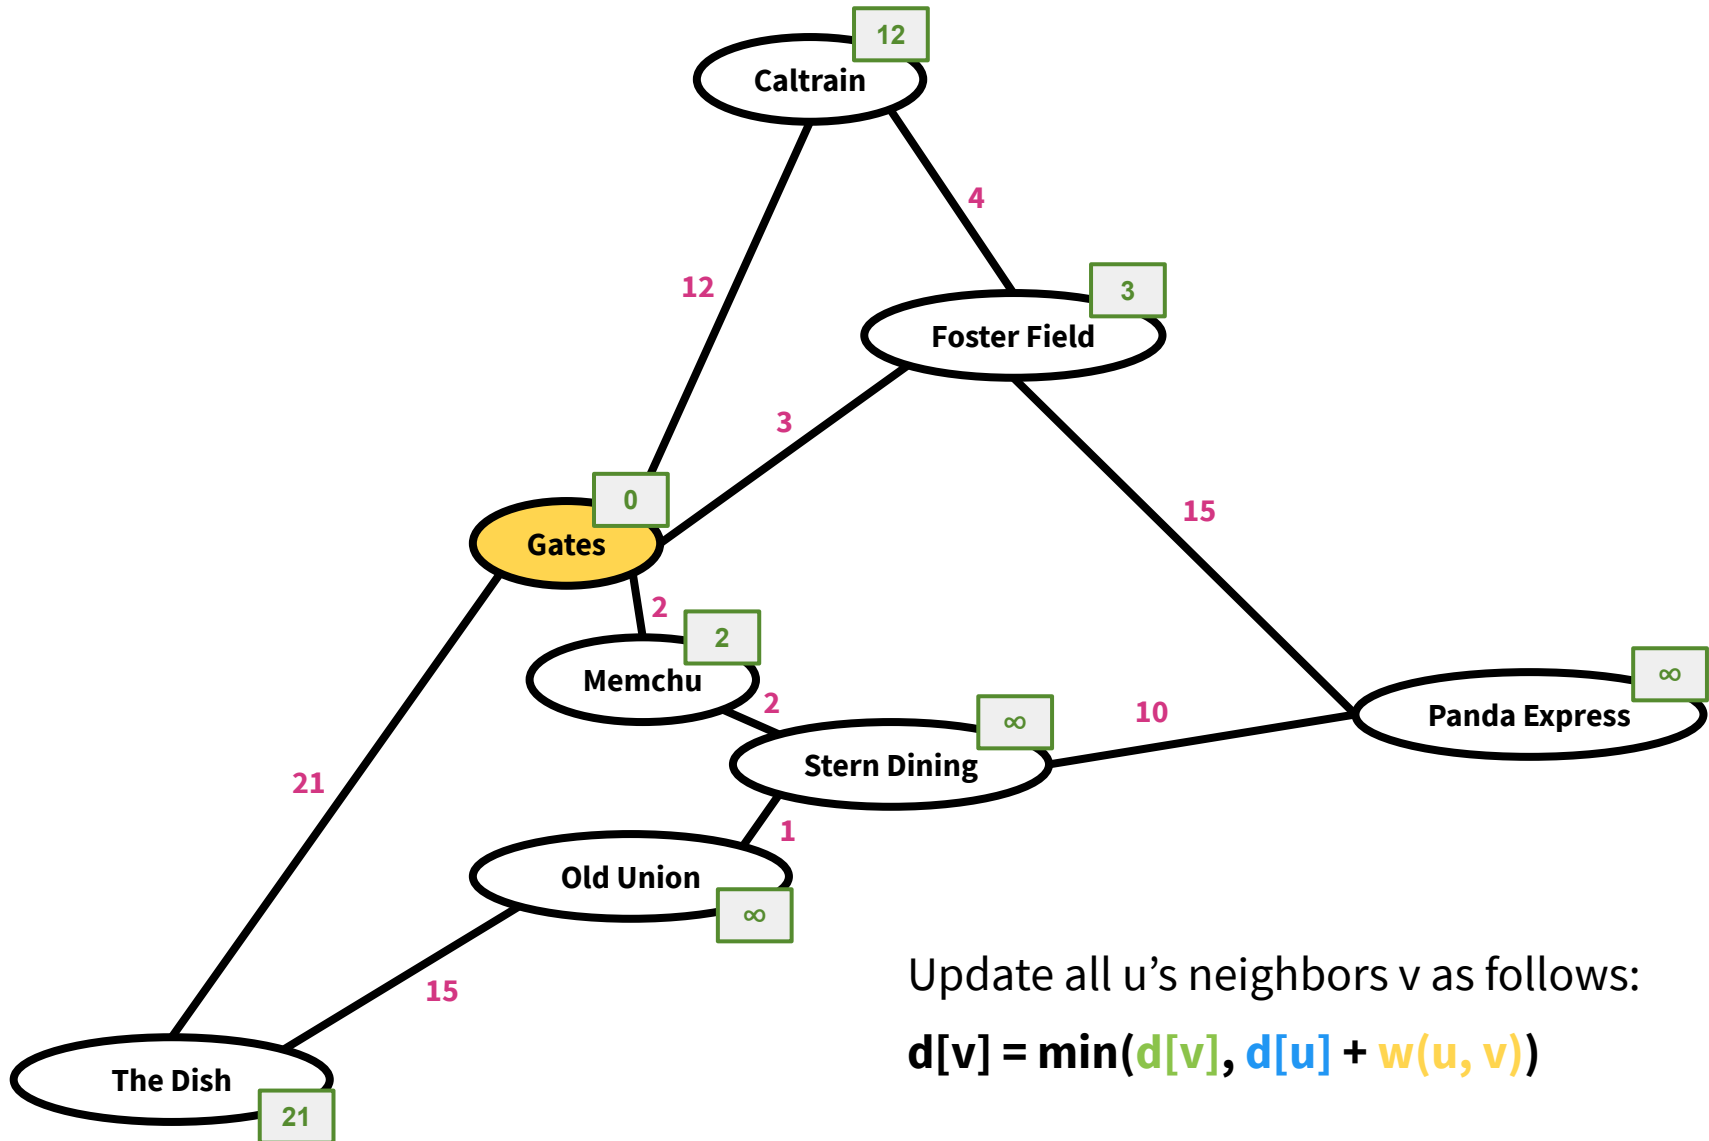
Dijkstra's Algorithm



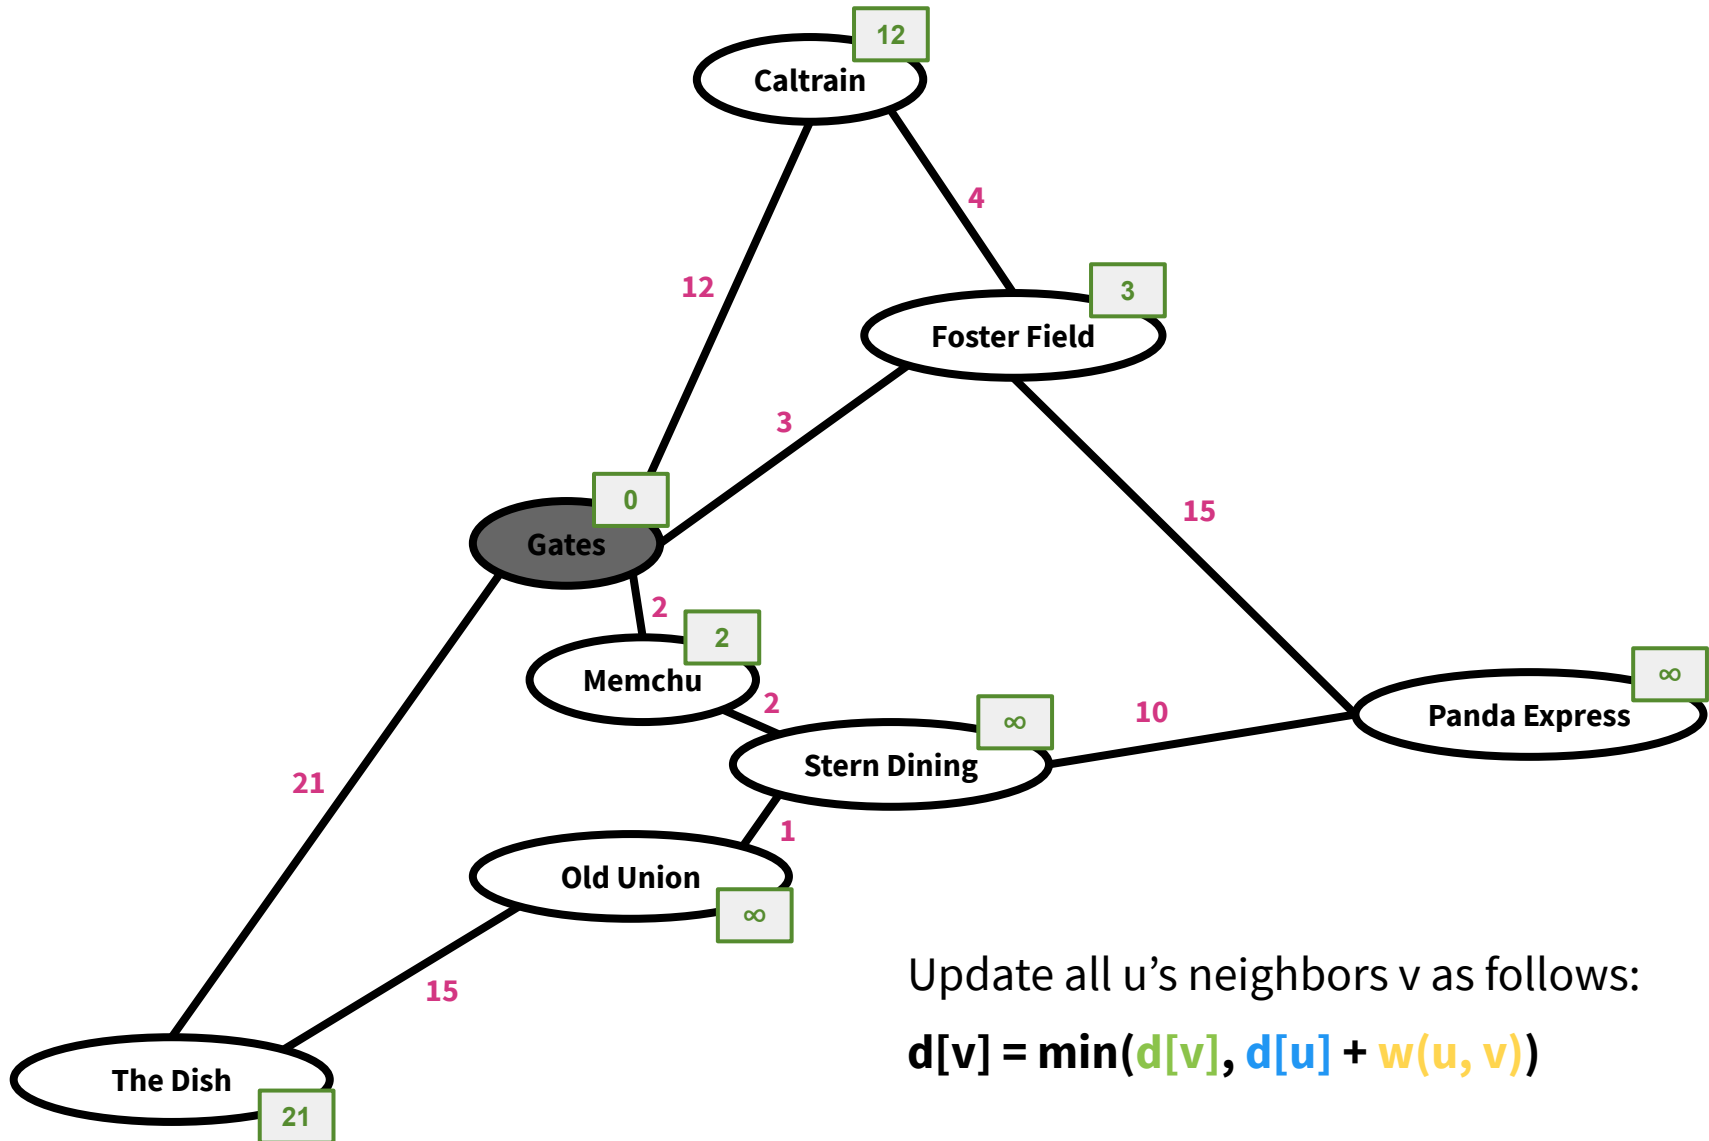
Update all u 's neighbors v as follows:

$$d[v] = \min(d[v], d[u] + w(u, v))$$

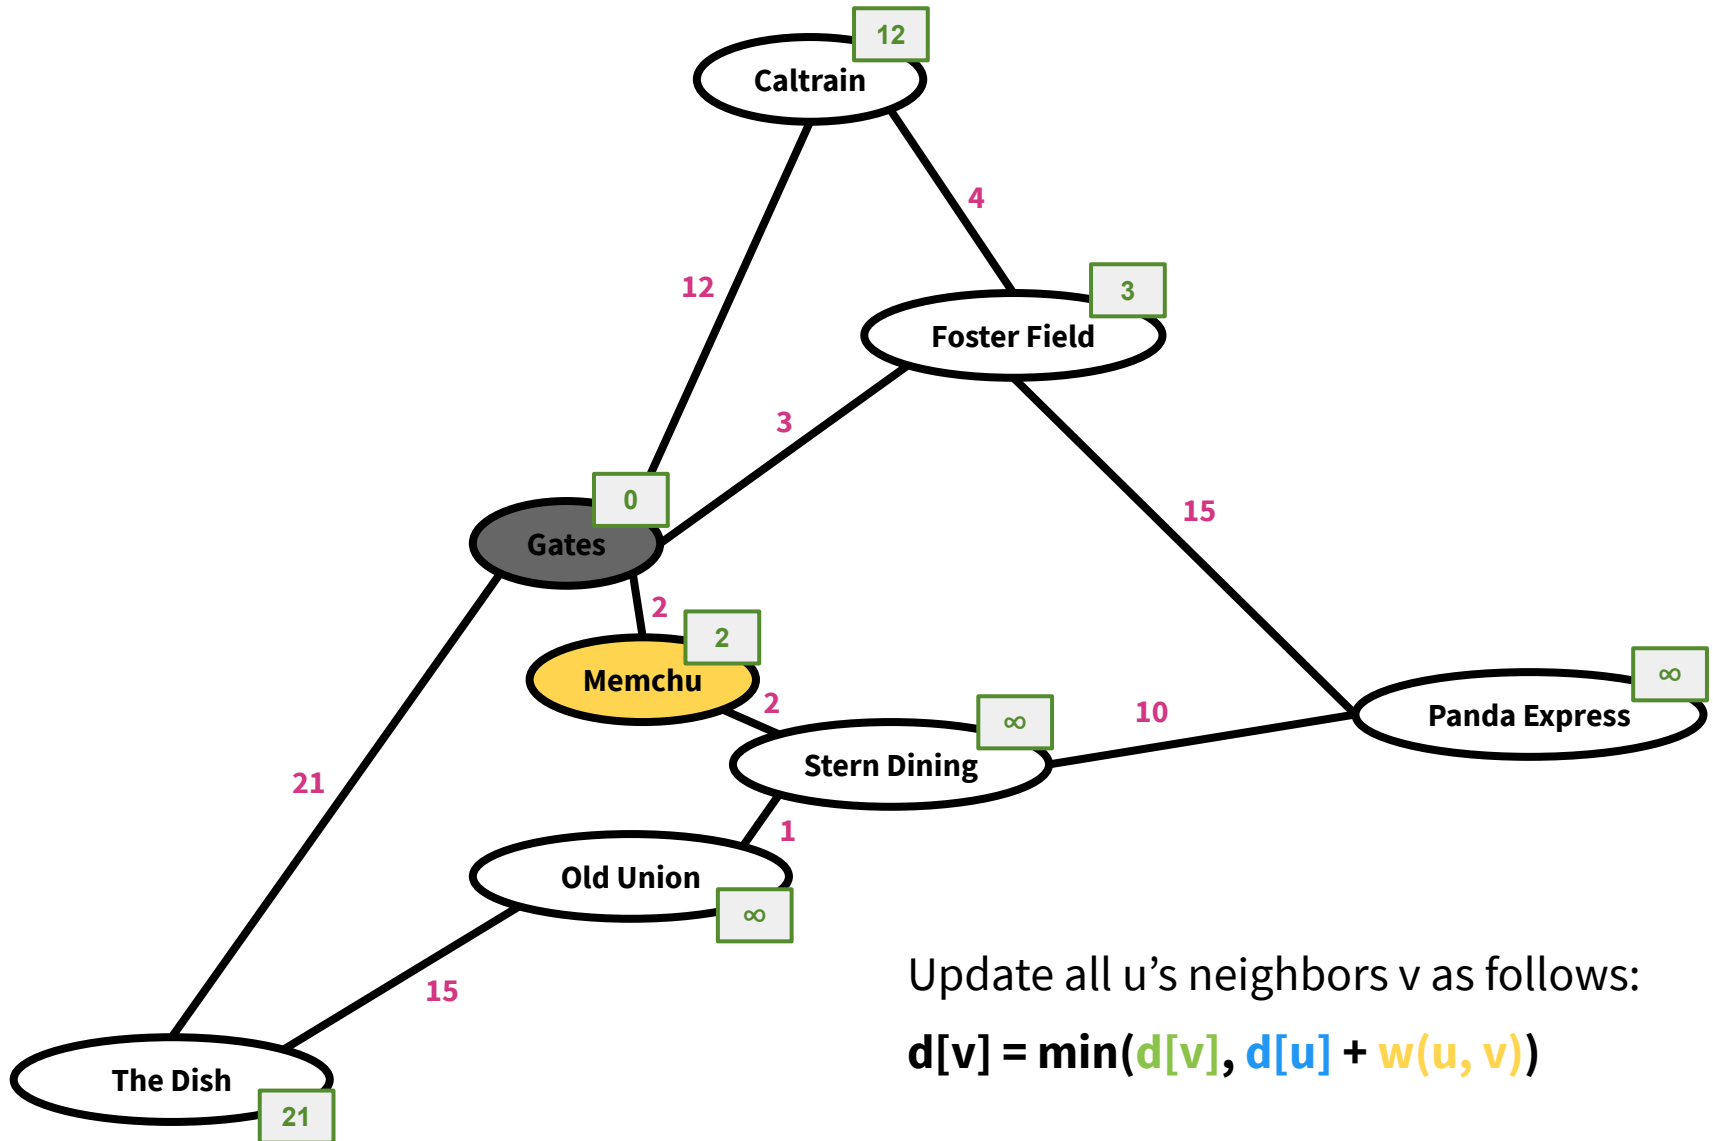
Dijkstra's Algorithm



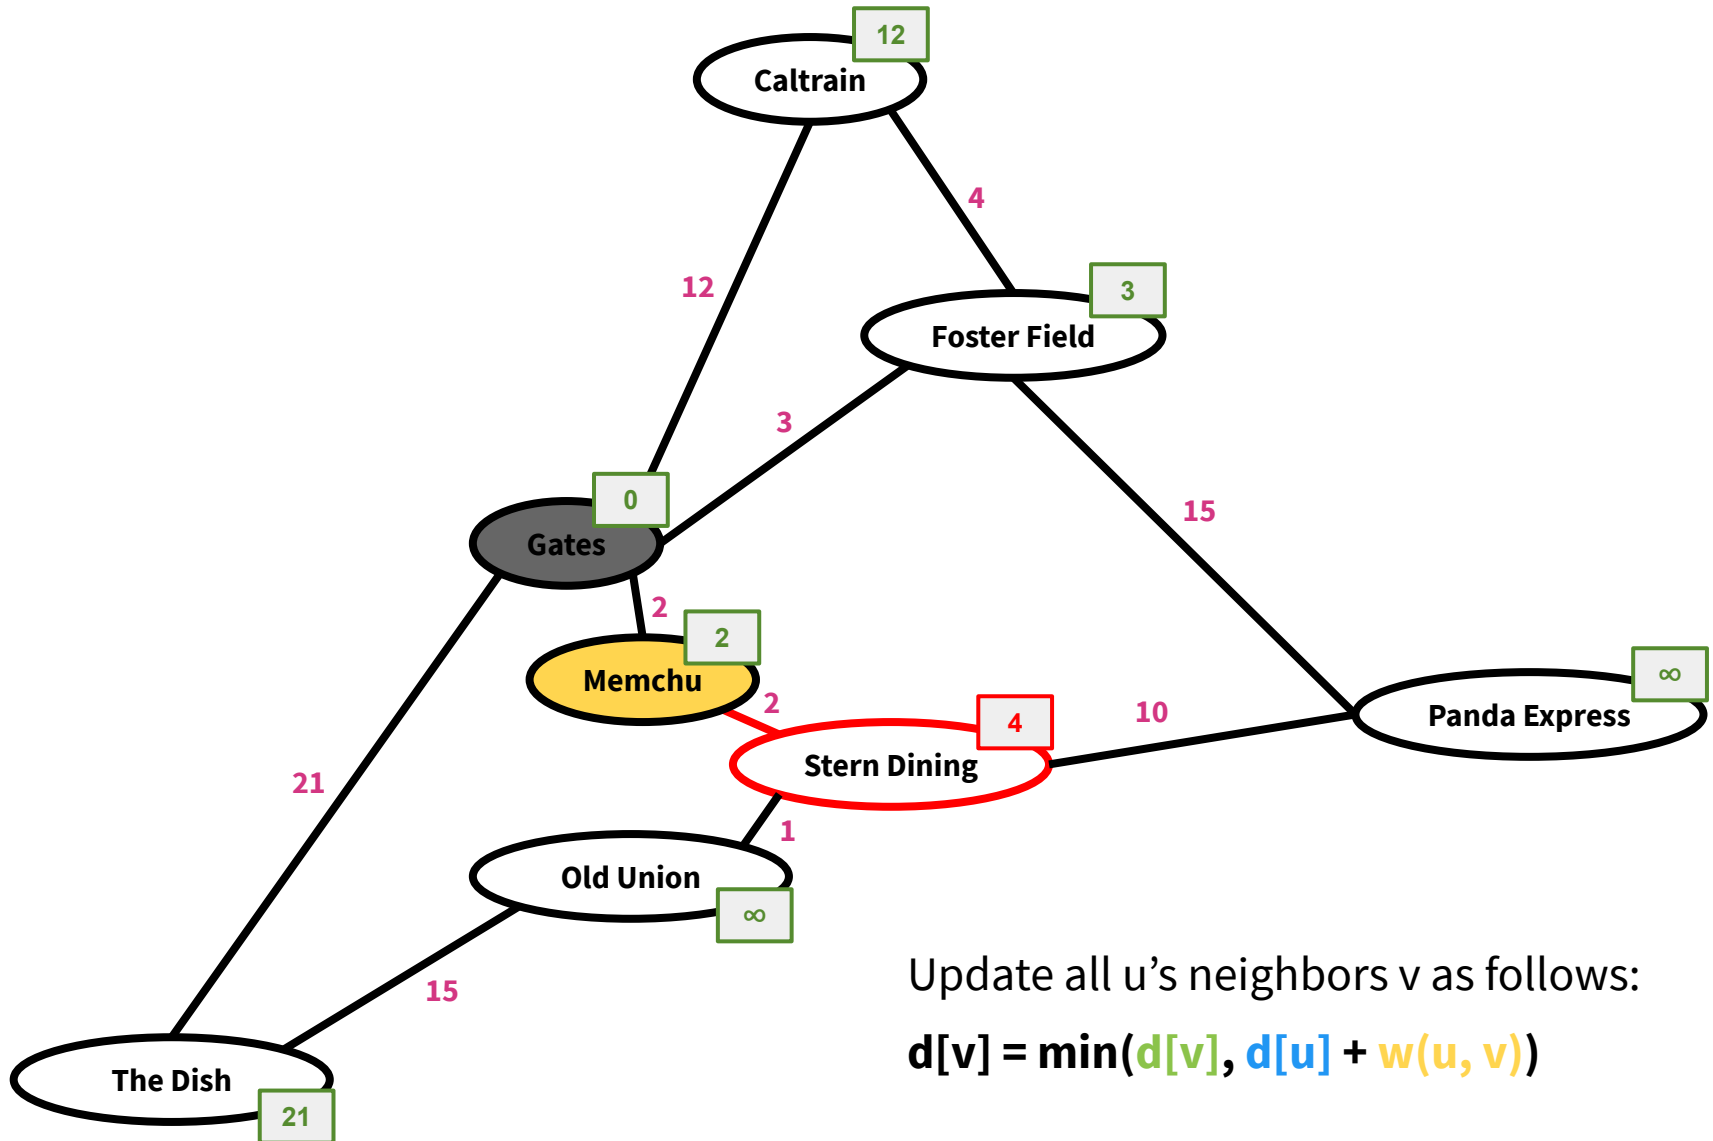
Dijkstra's Algorithm



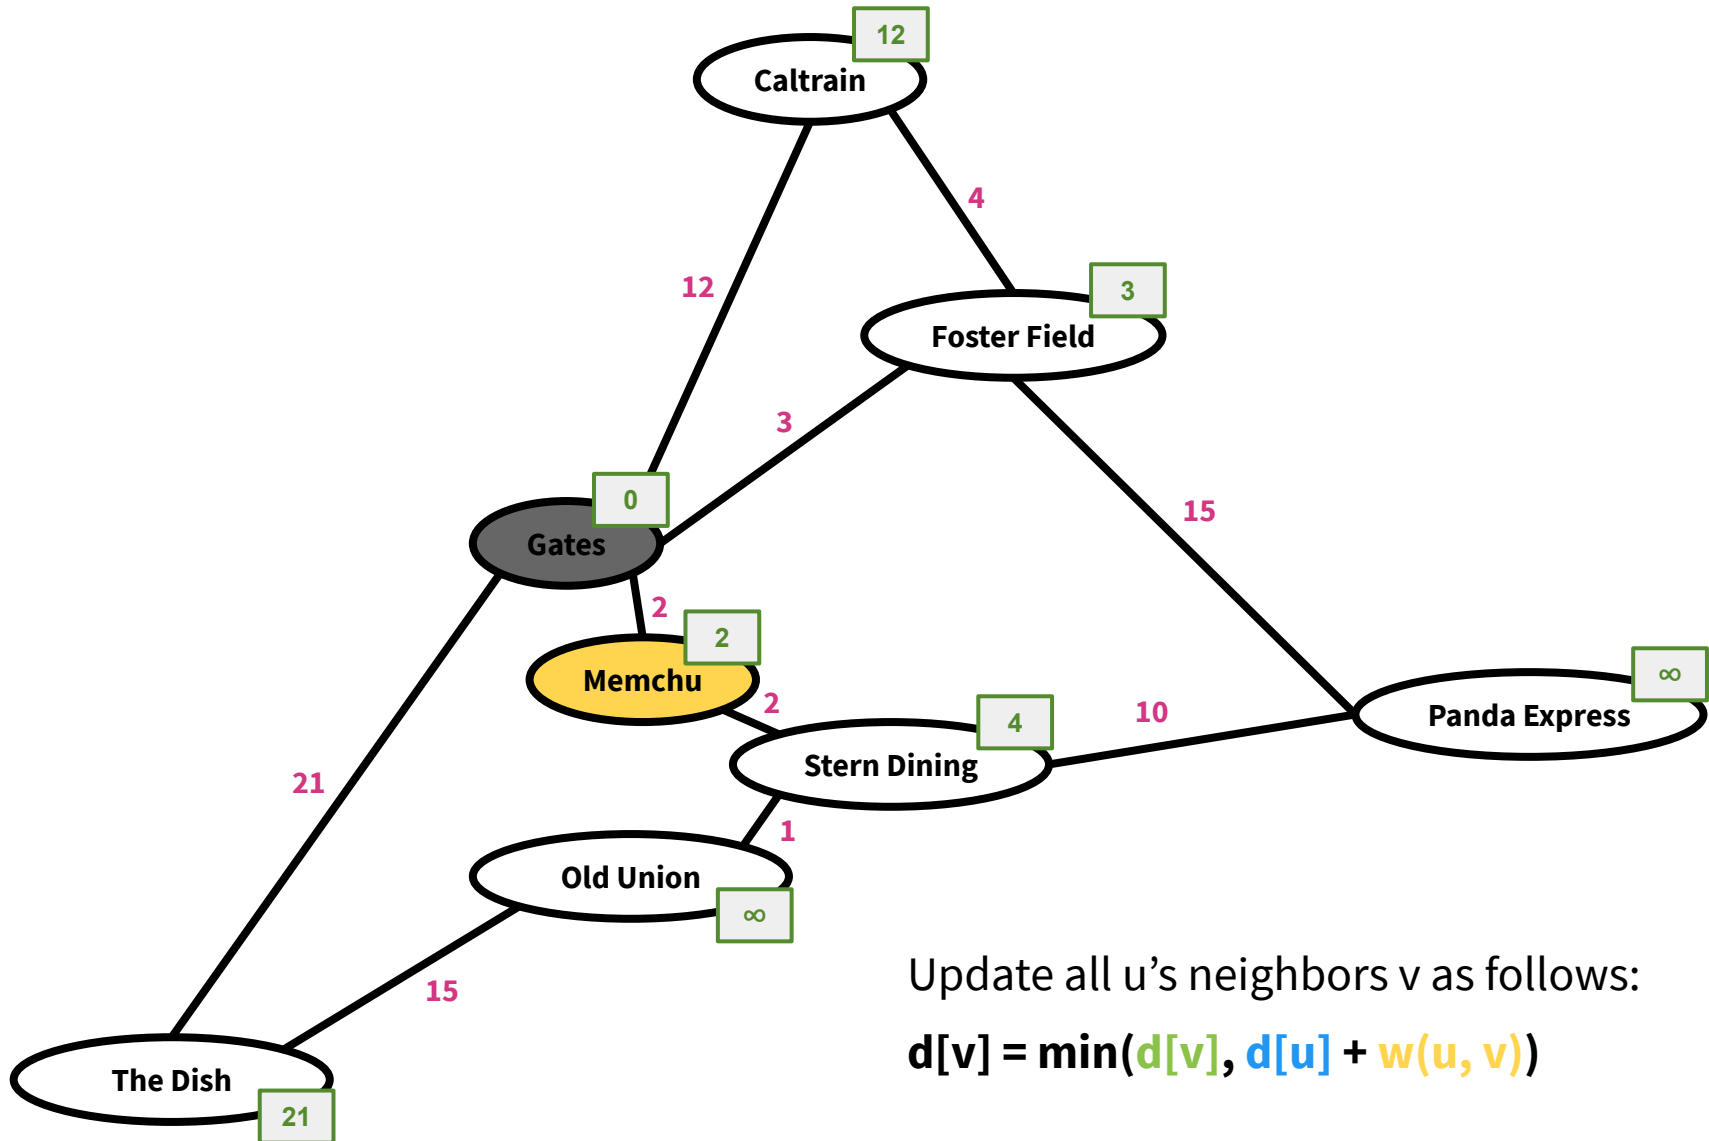
Dijkstra's Algorithm



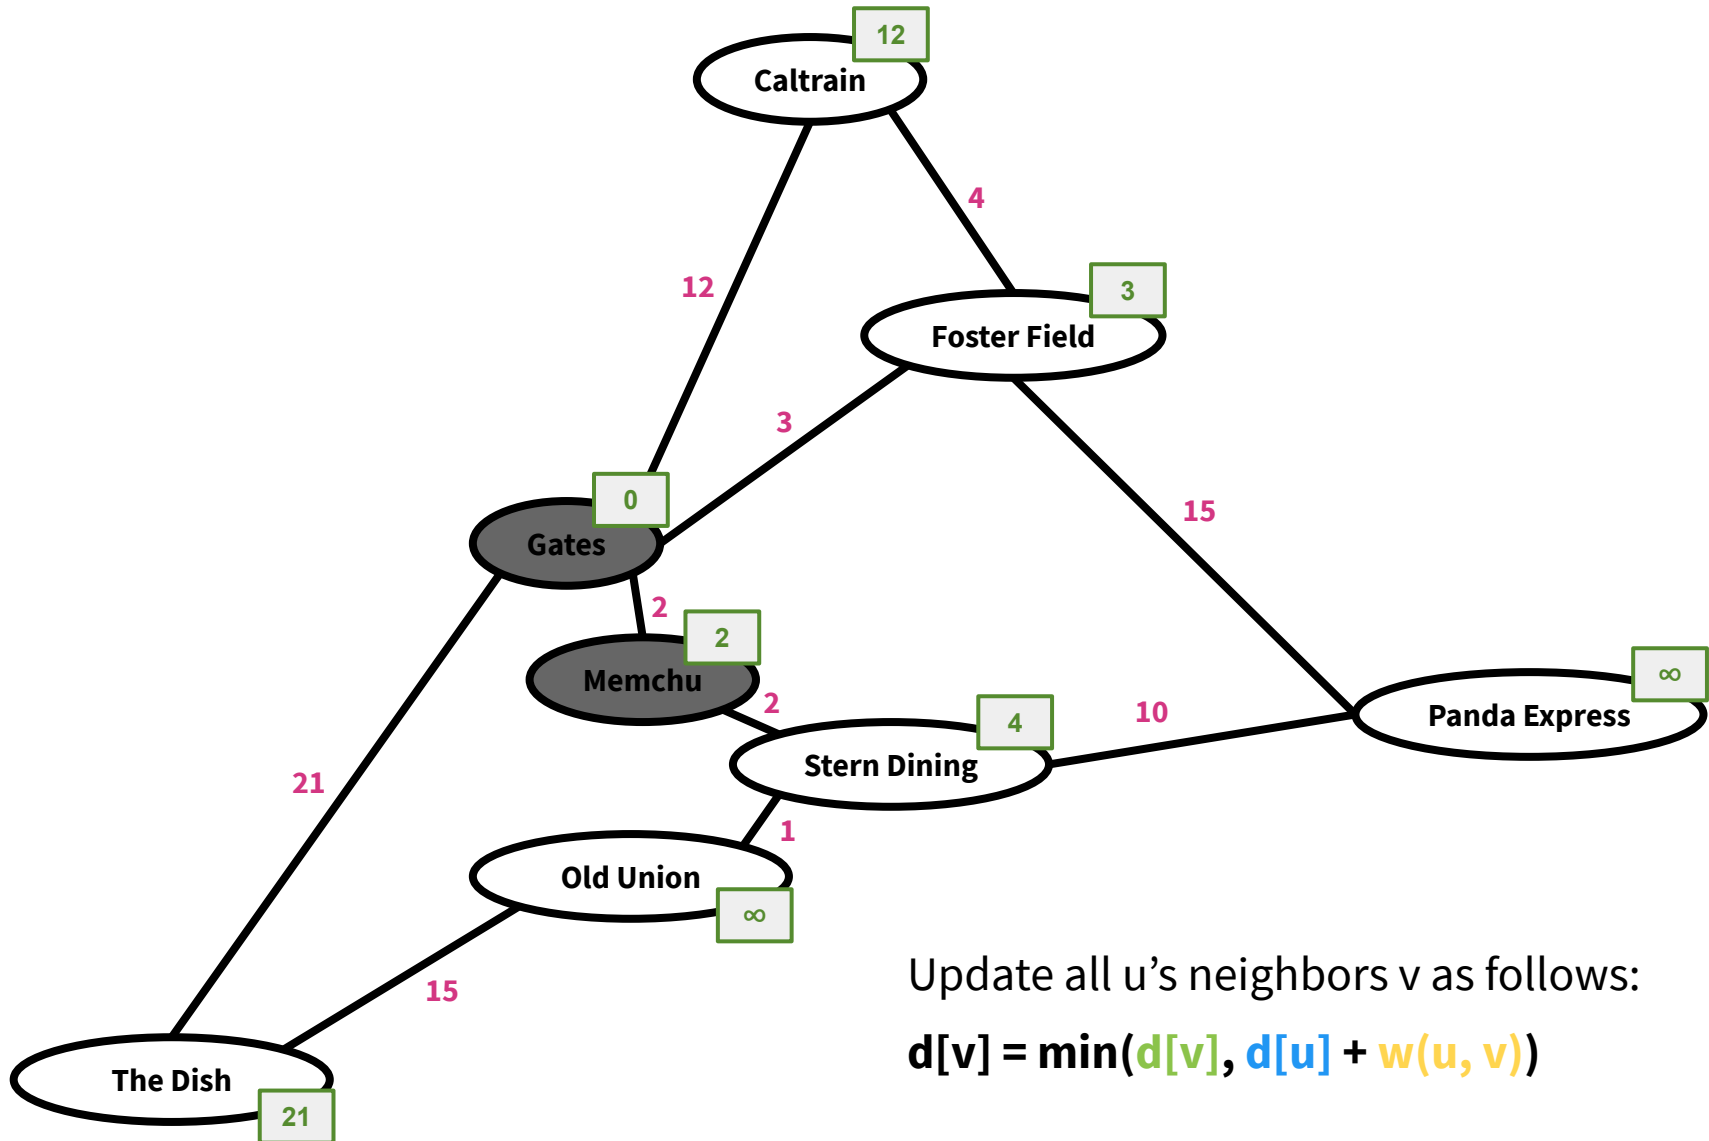
Dijkstra's Algorithm



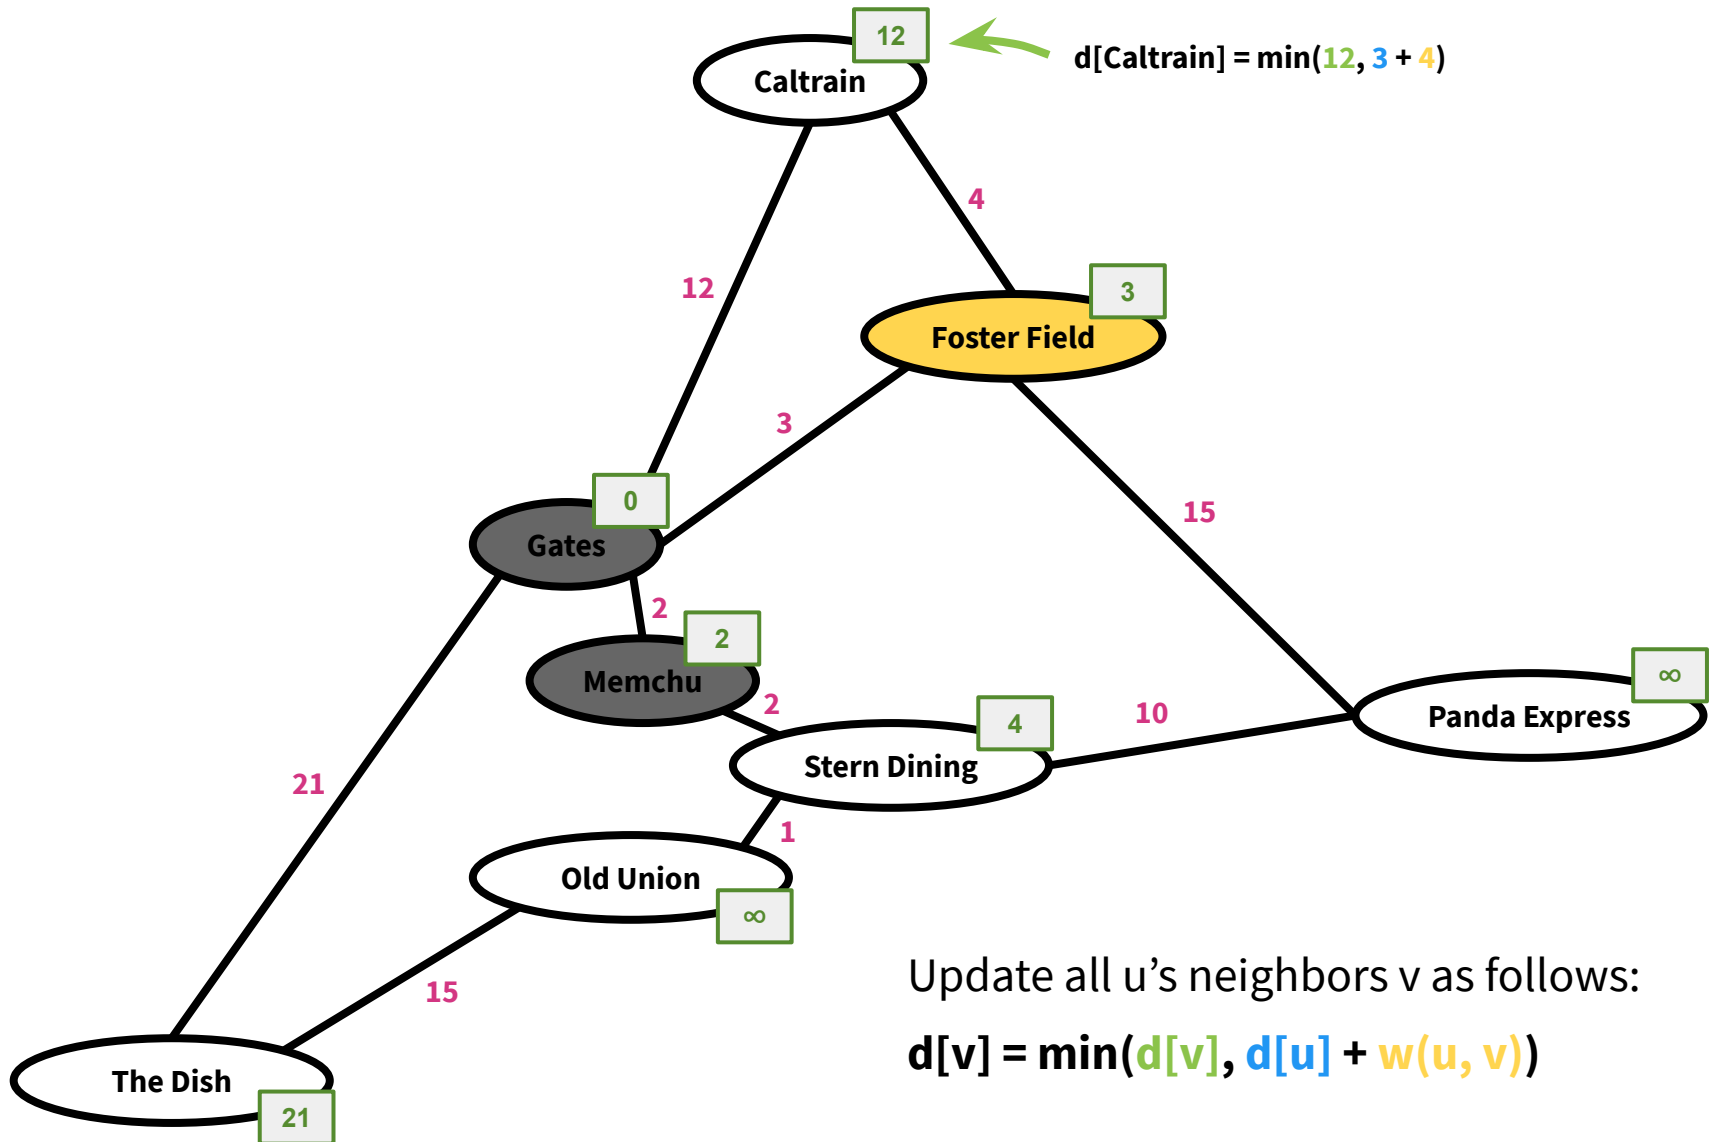
Dijkstra's Algorithm



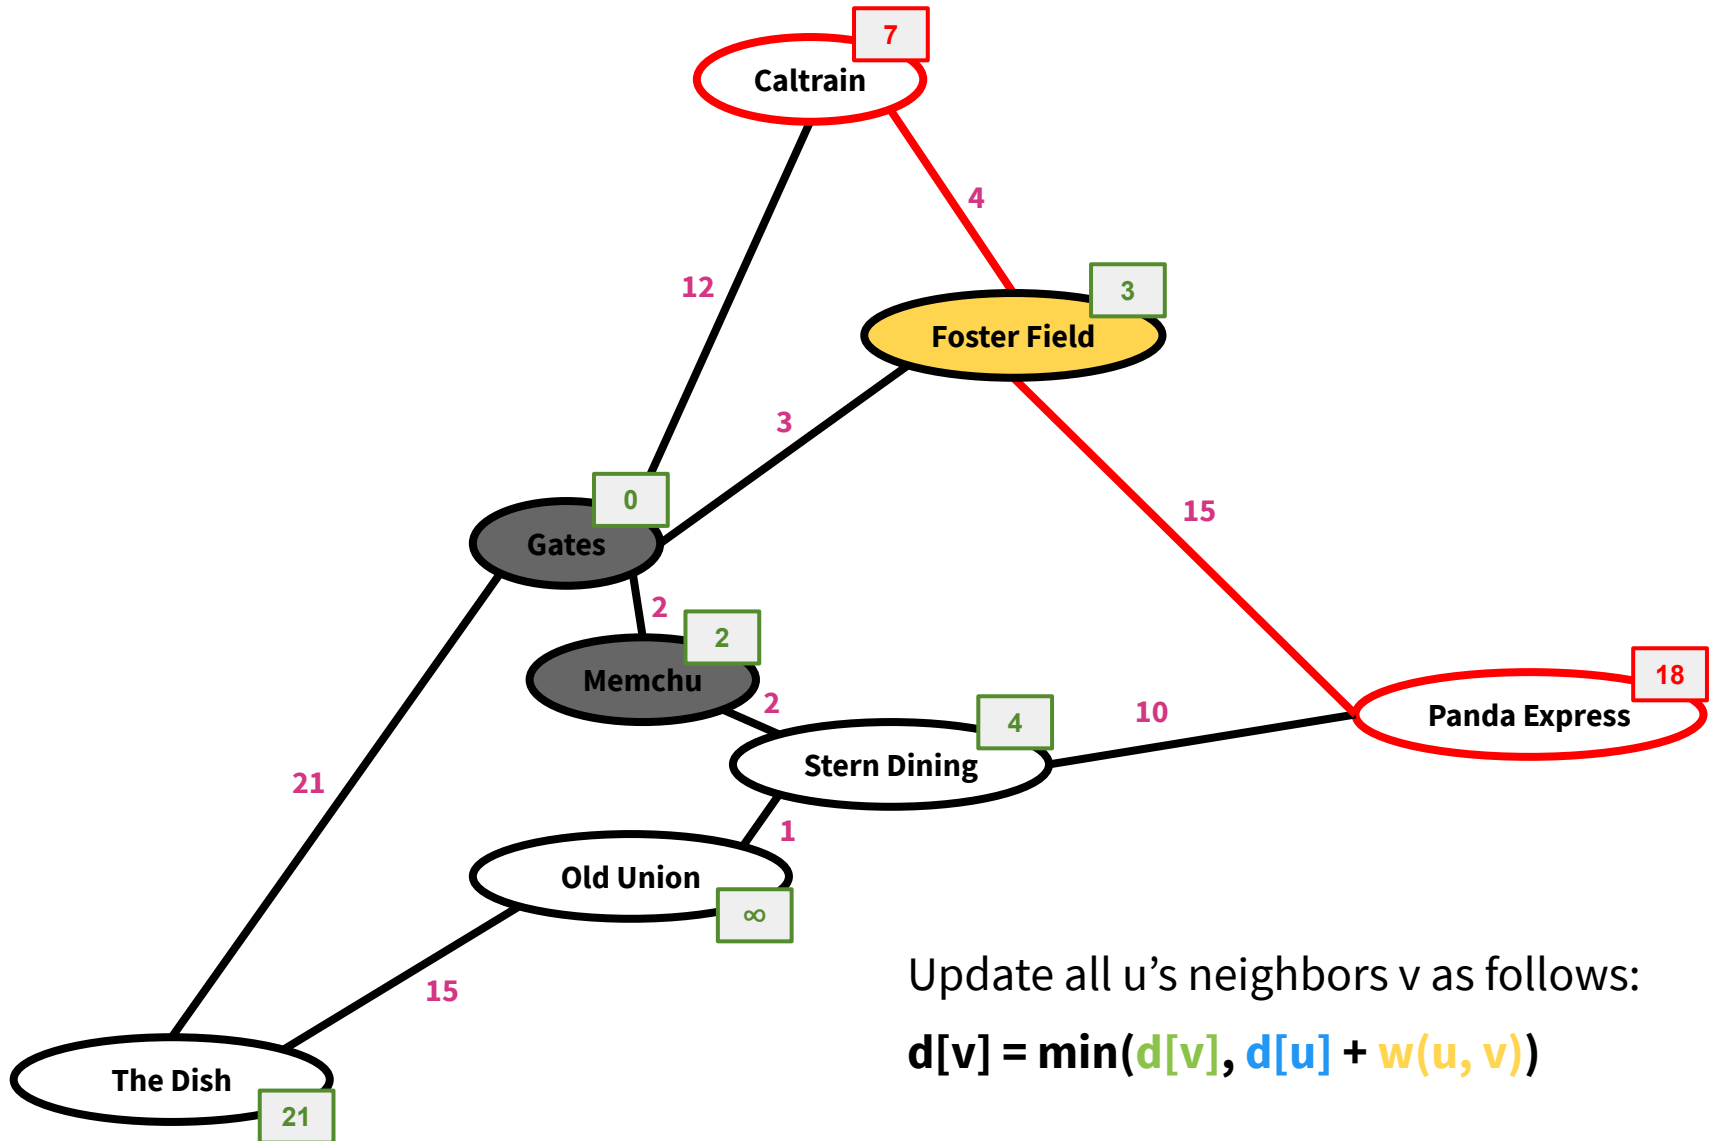
Dijkstra's Algorithm



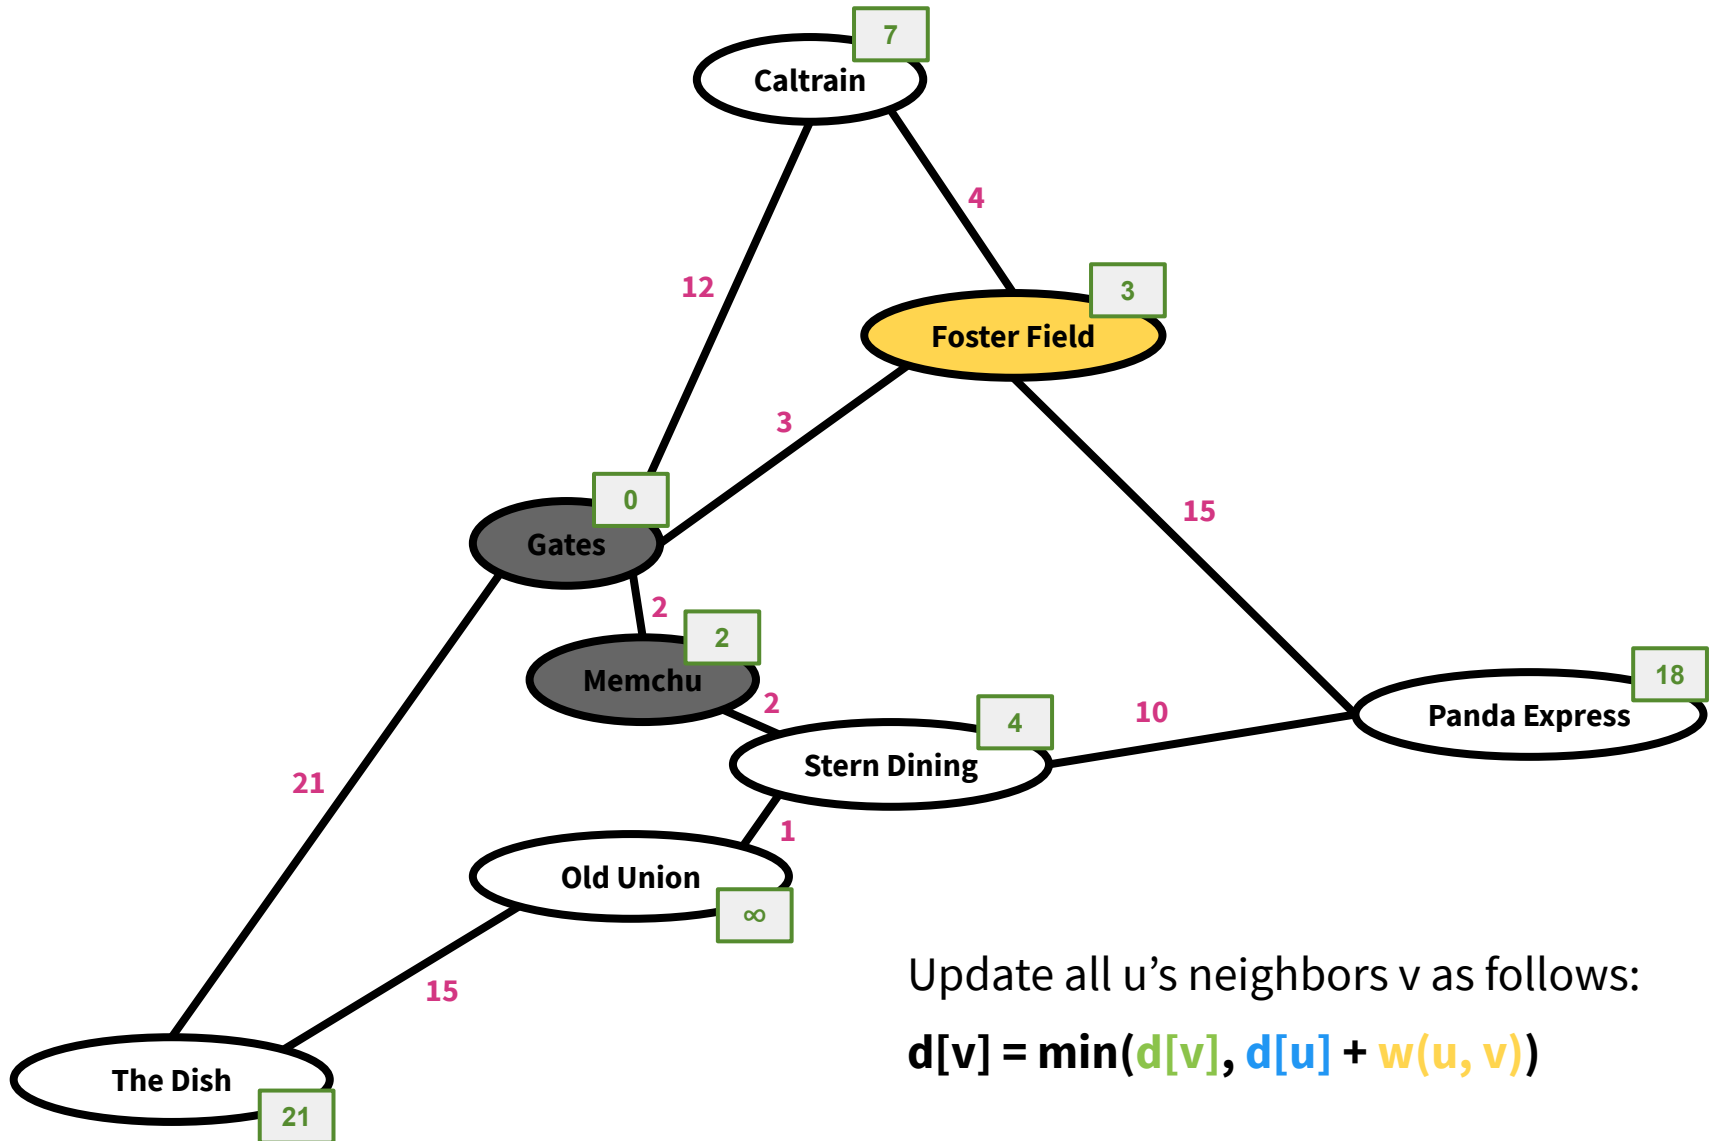
Dijkstra's Algorithm



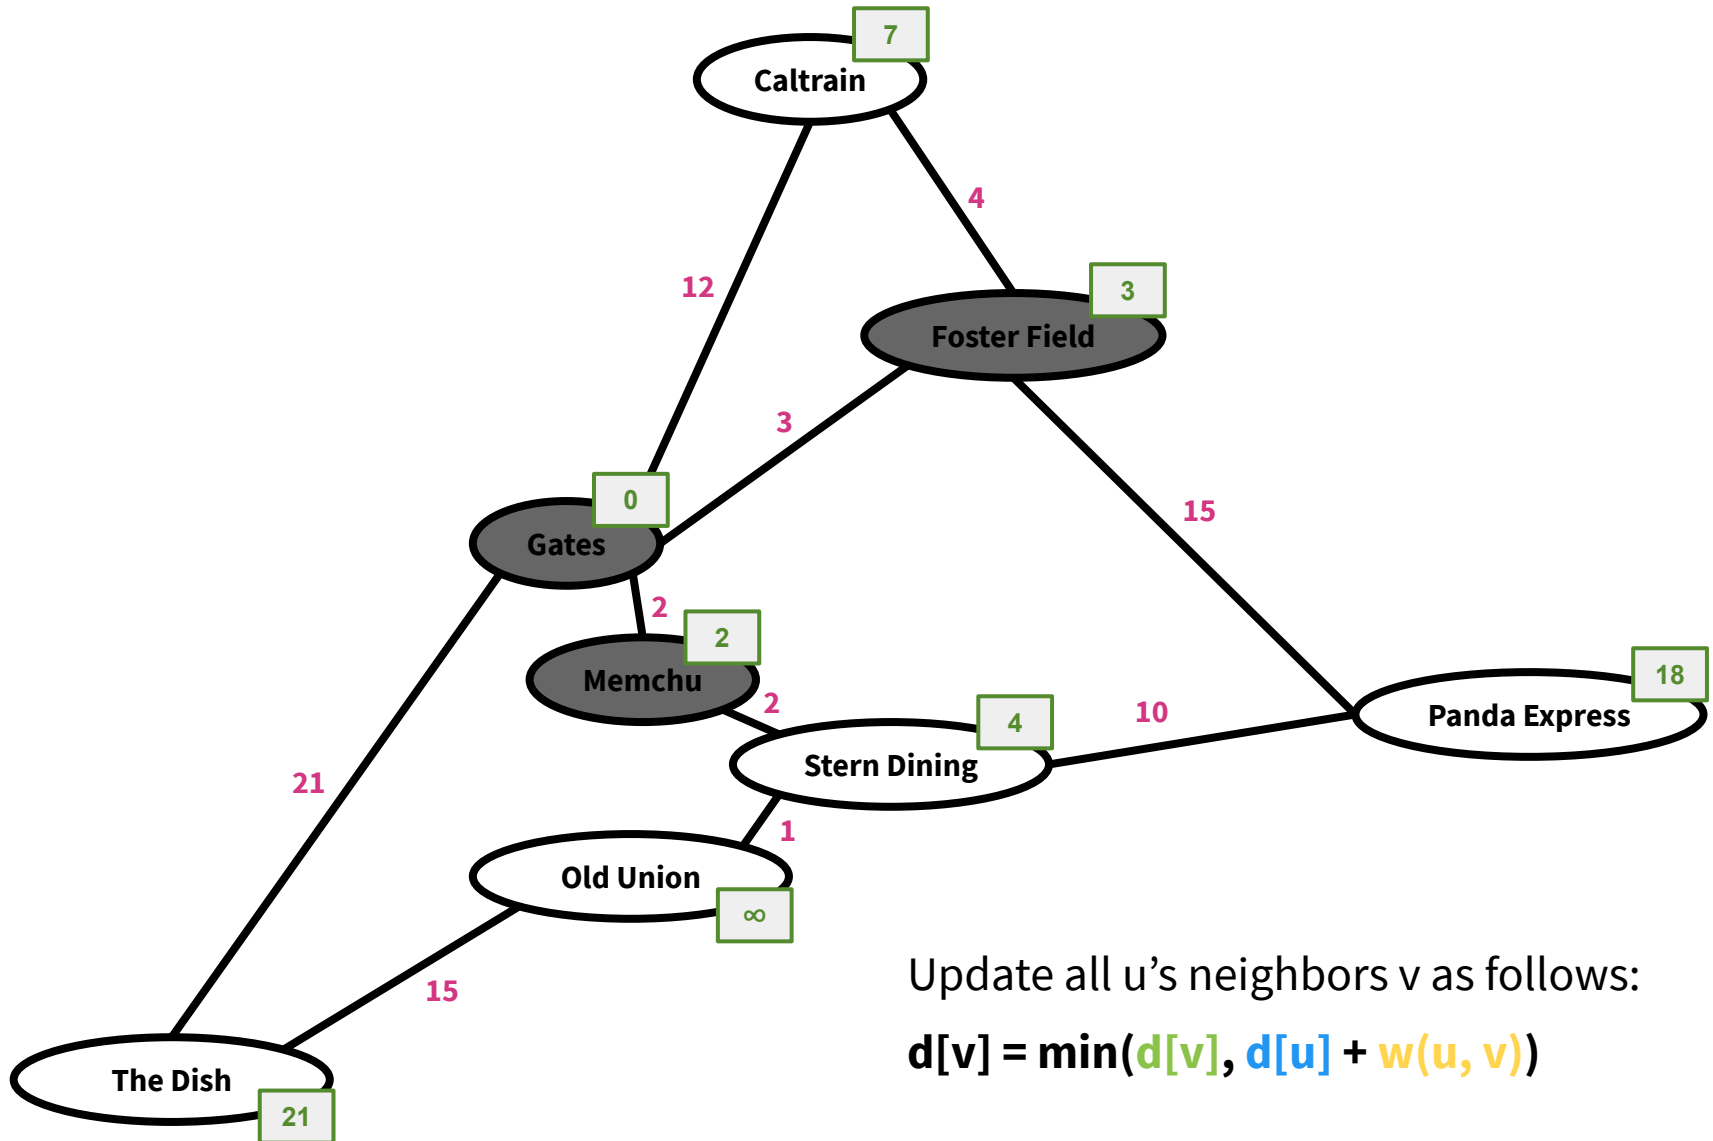
Dijkstra's Algorithm



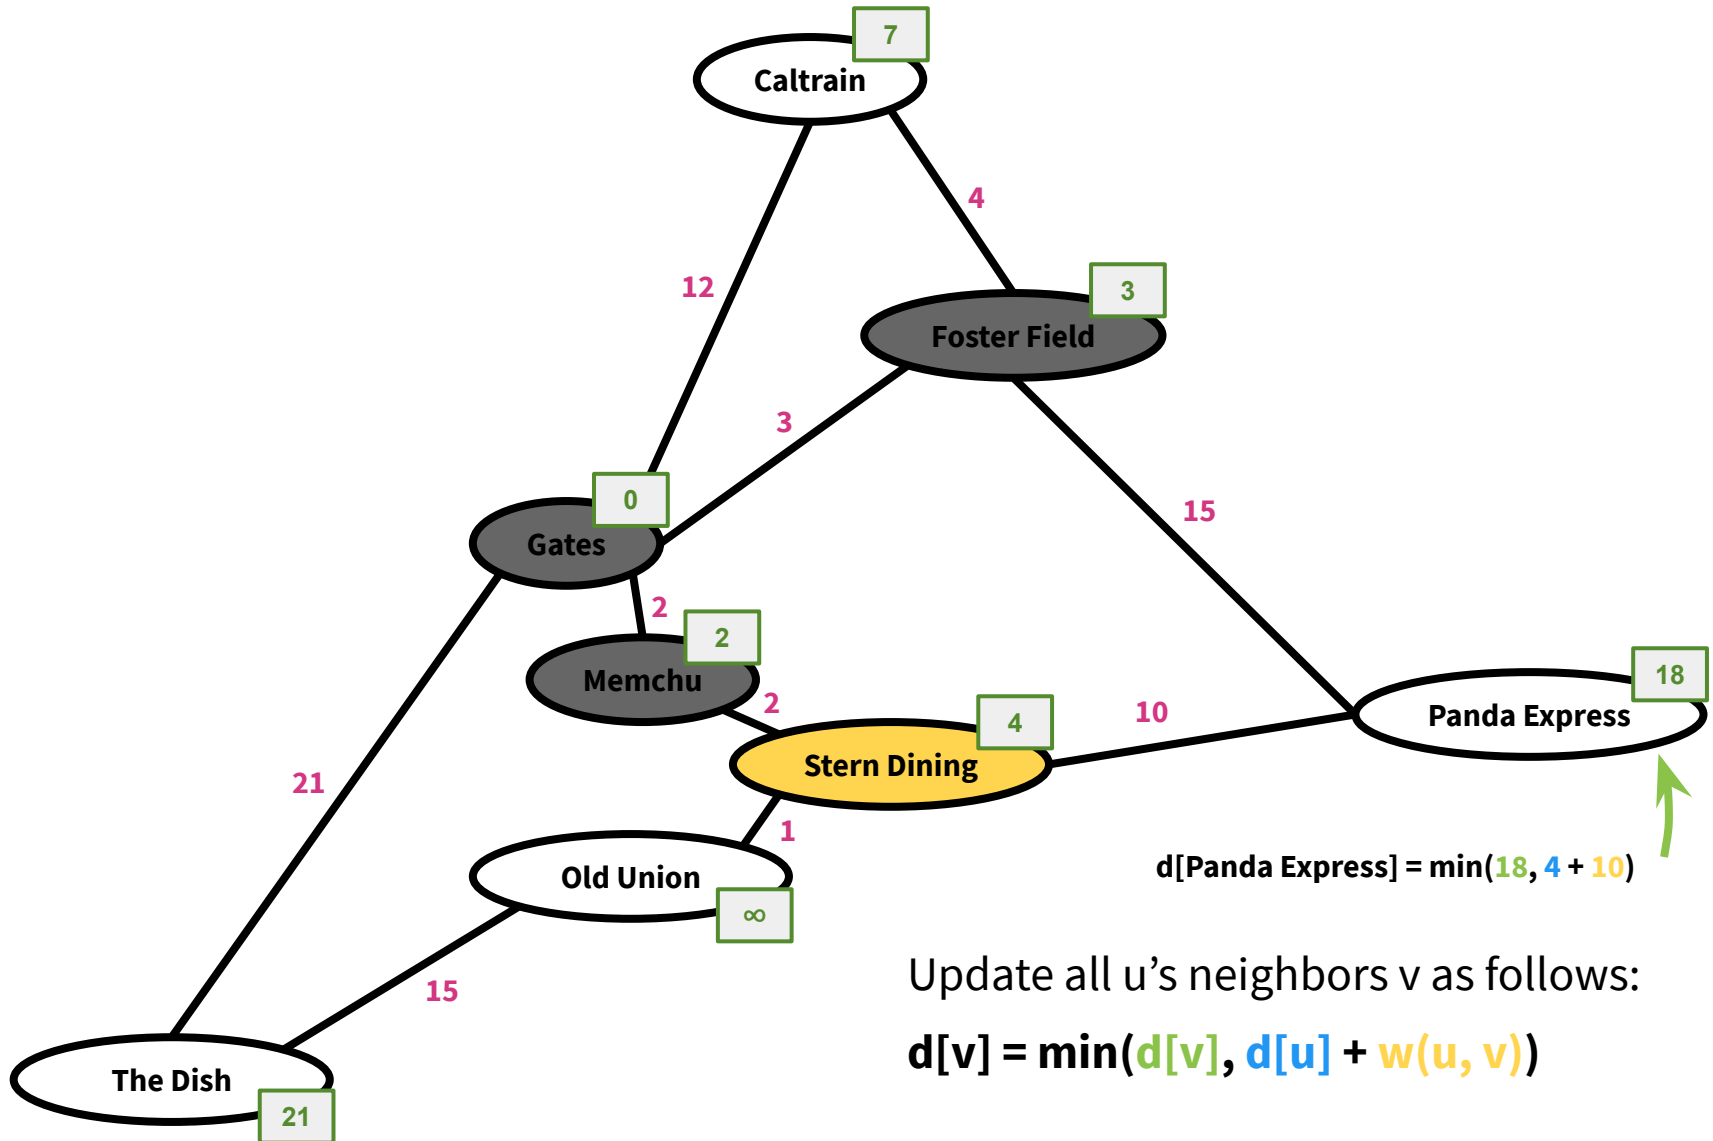
Dijkstra's Algorithm



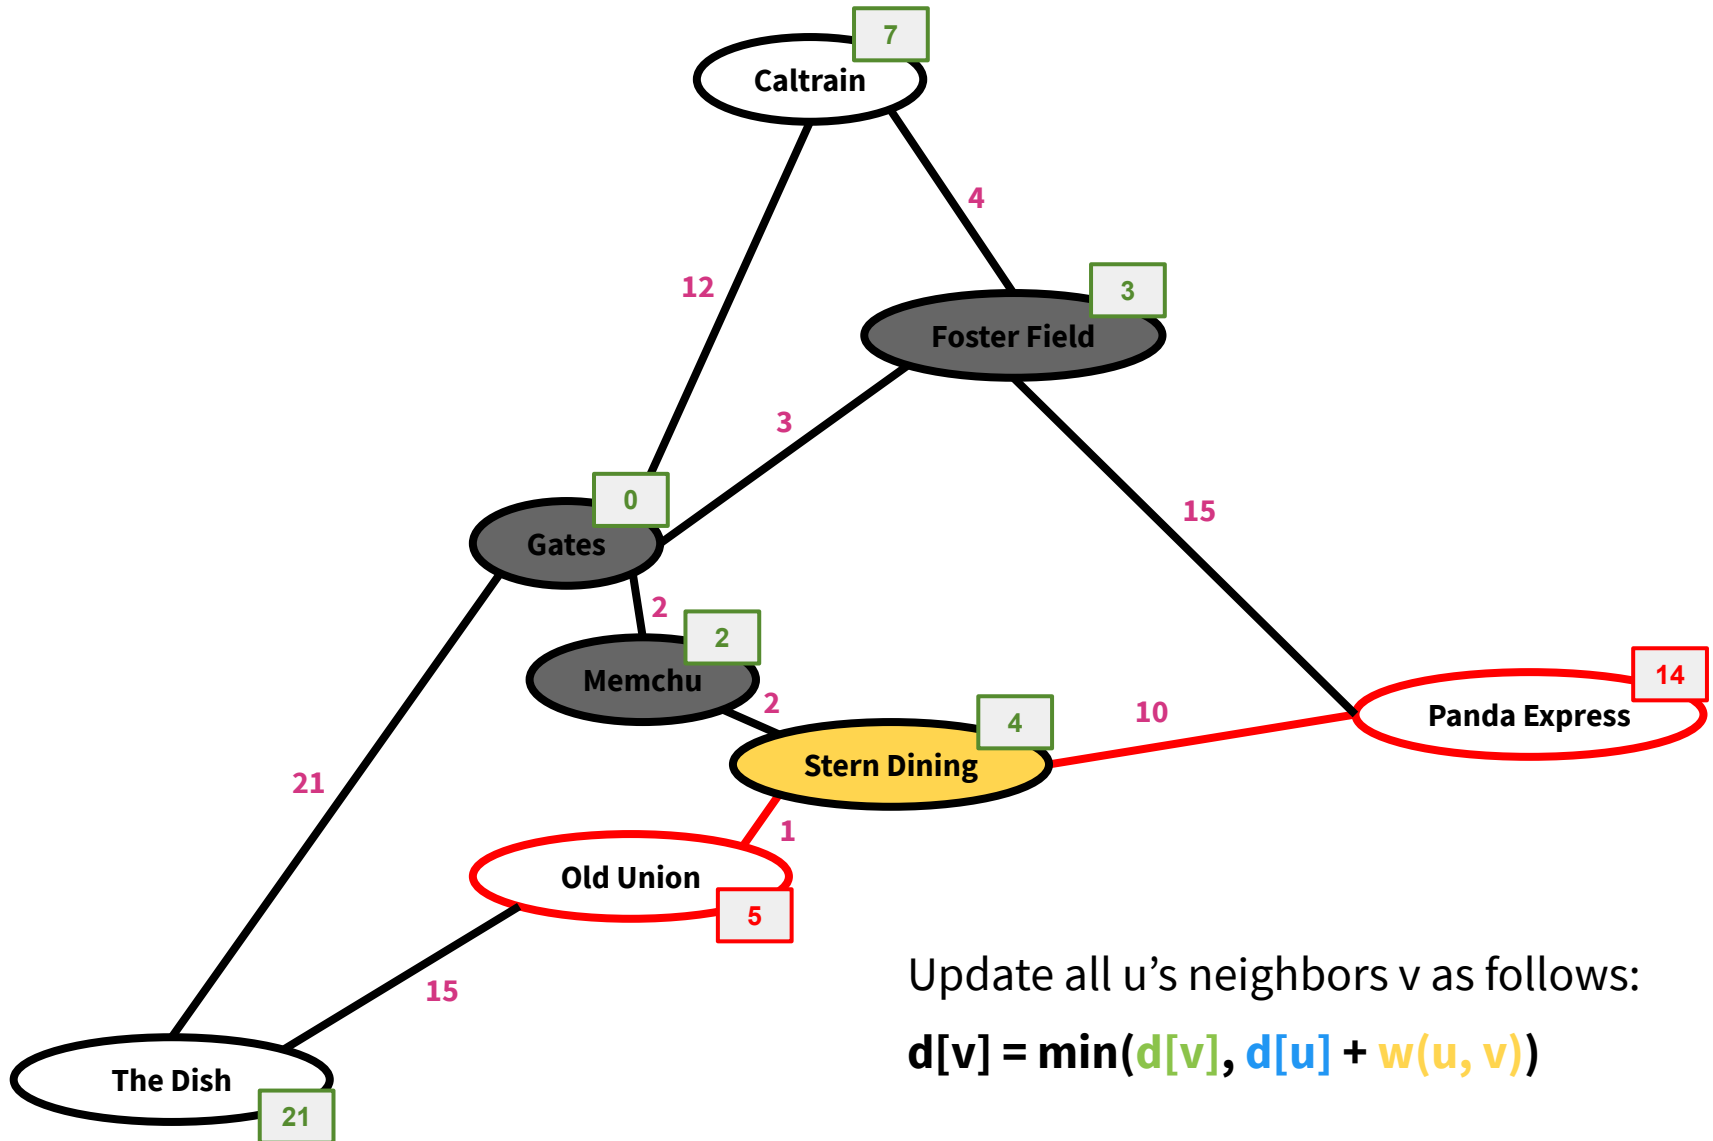
Dijkstra's Algorithm



Dijkstra's Algorithm



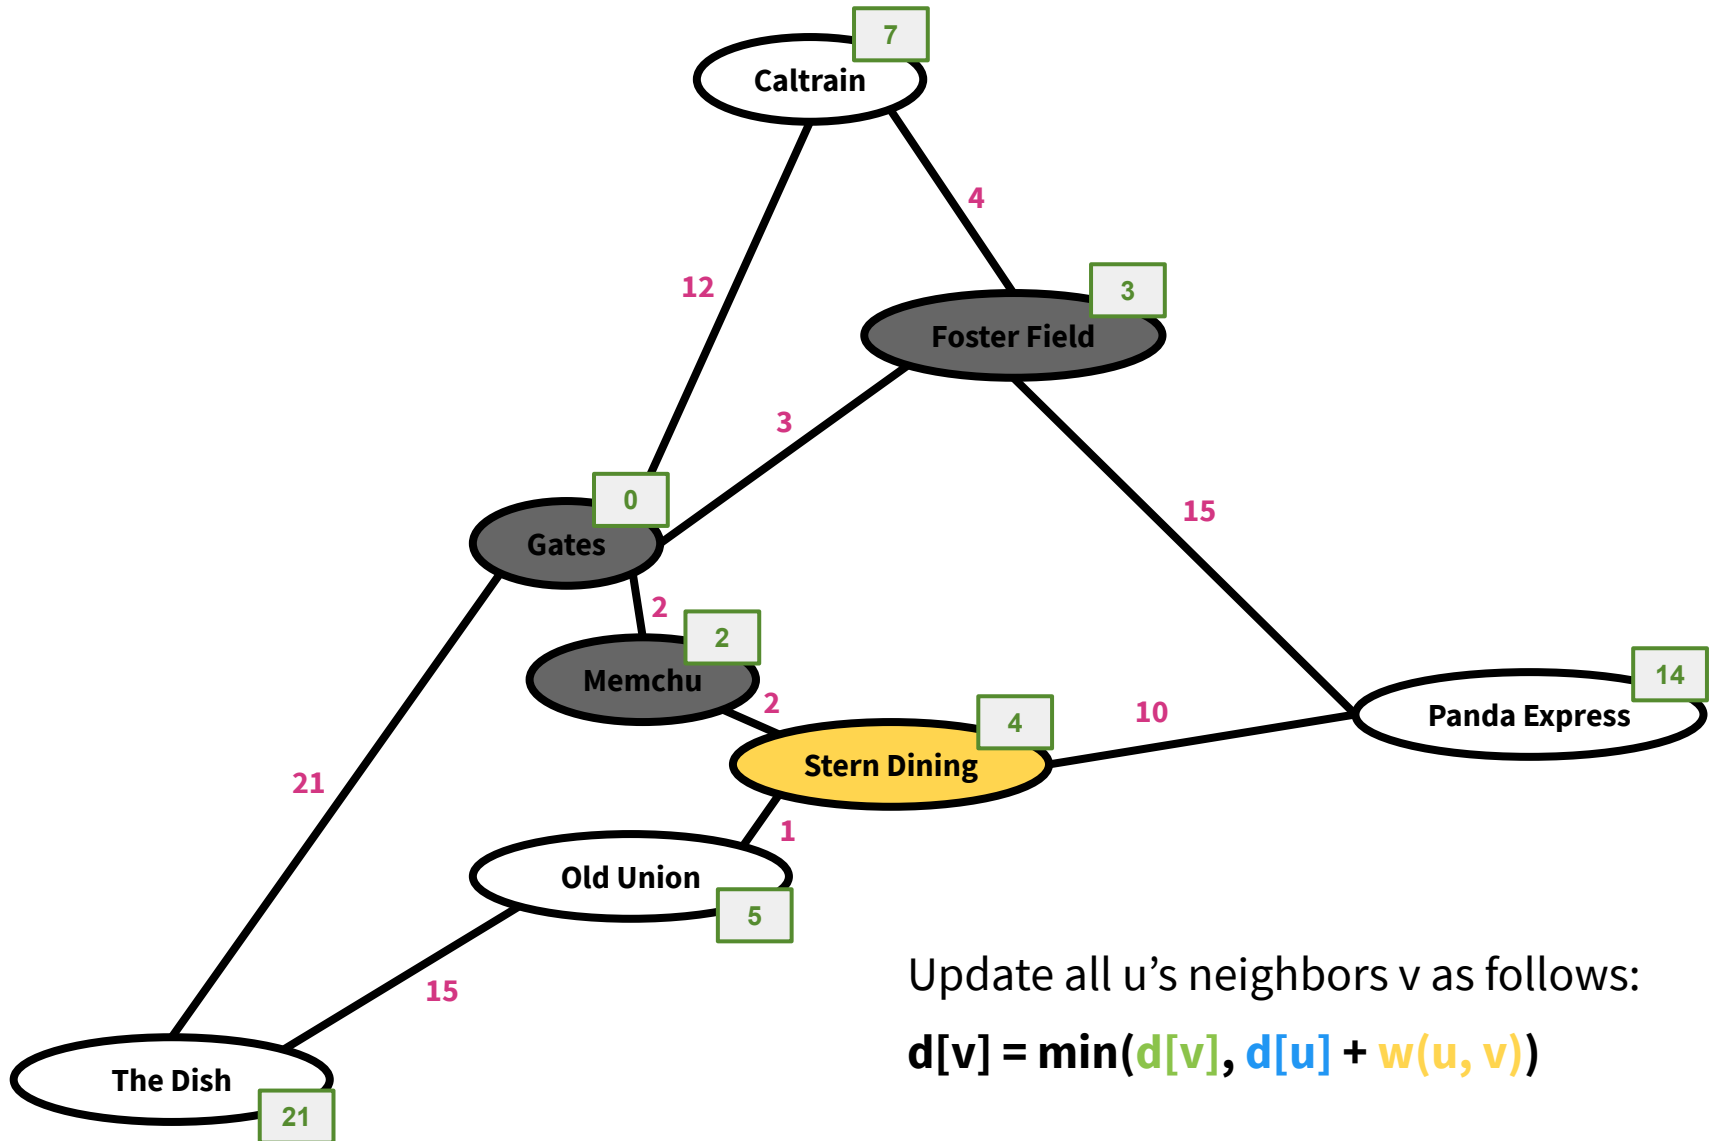
Dijkstra's Algorithm



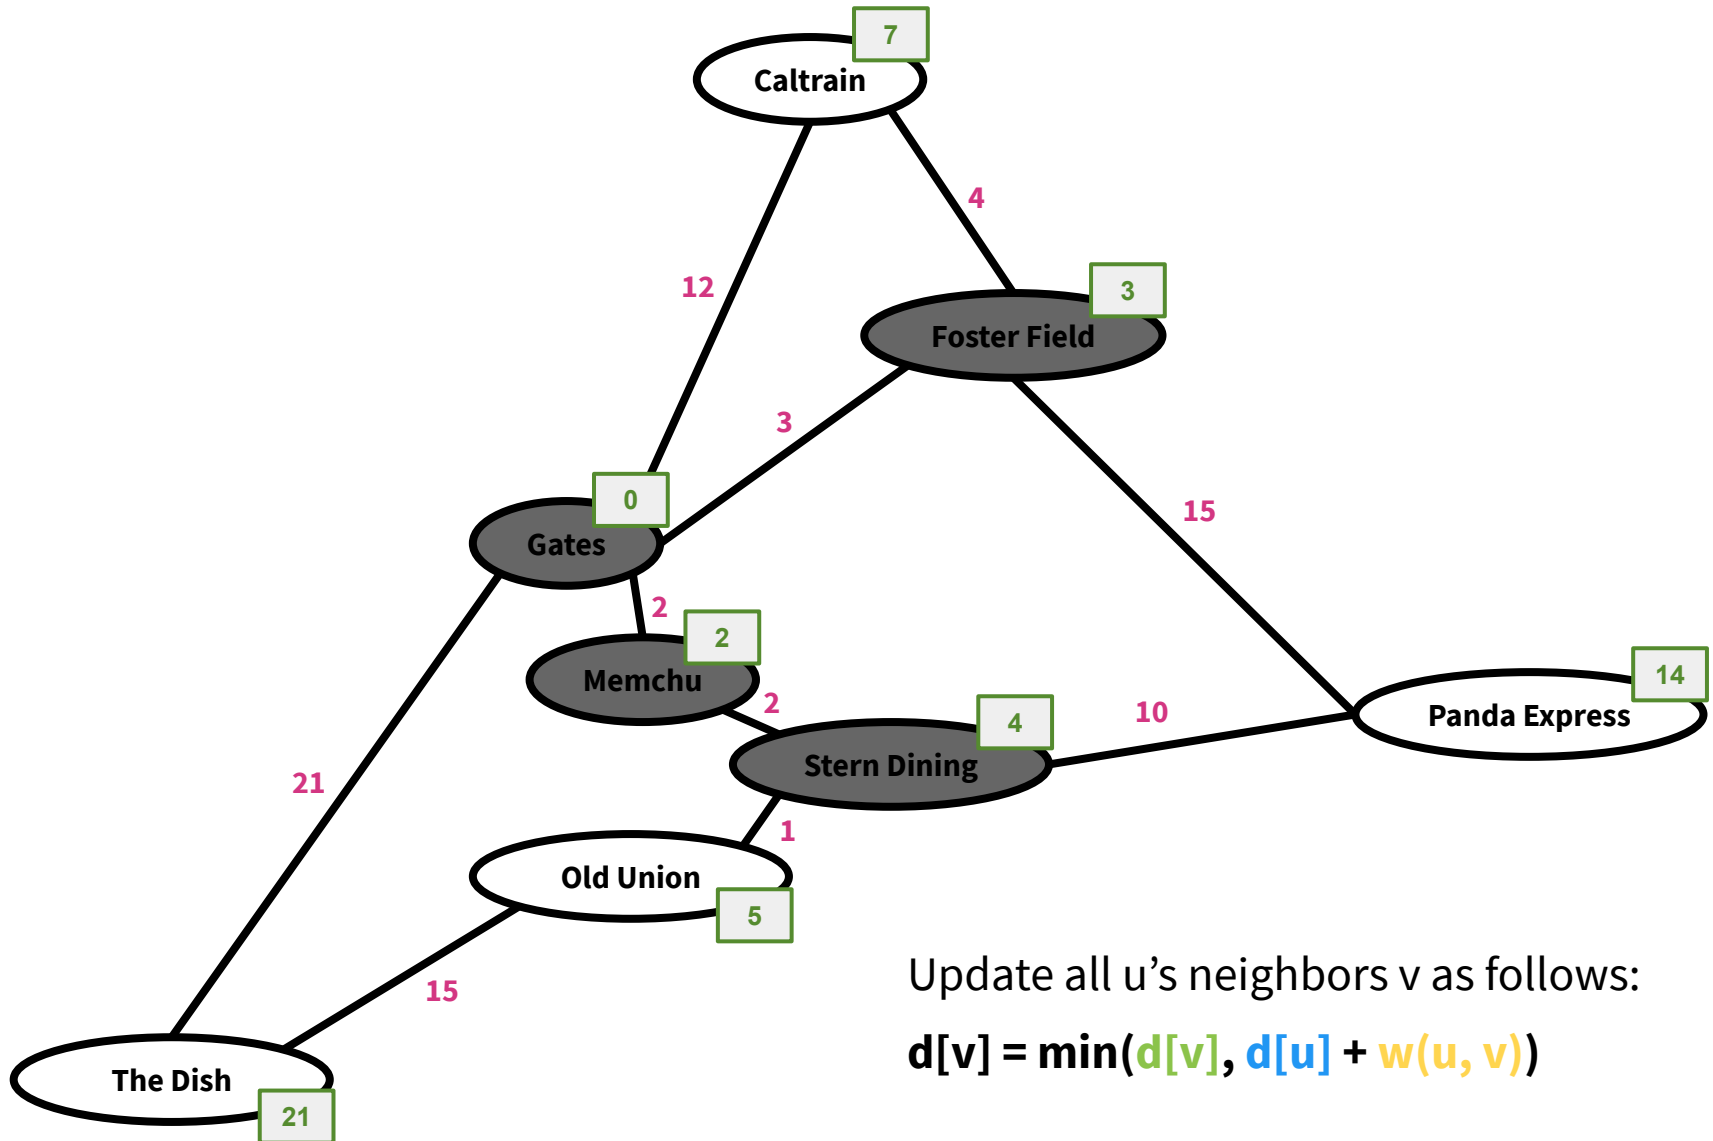
Update all u 's neighbors v as follows:

$$d[v] = \min(d[v], d[u] + w(u, v))$$

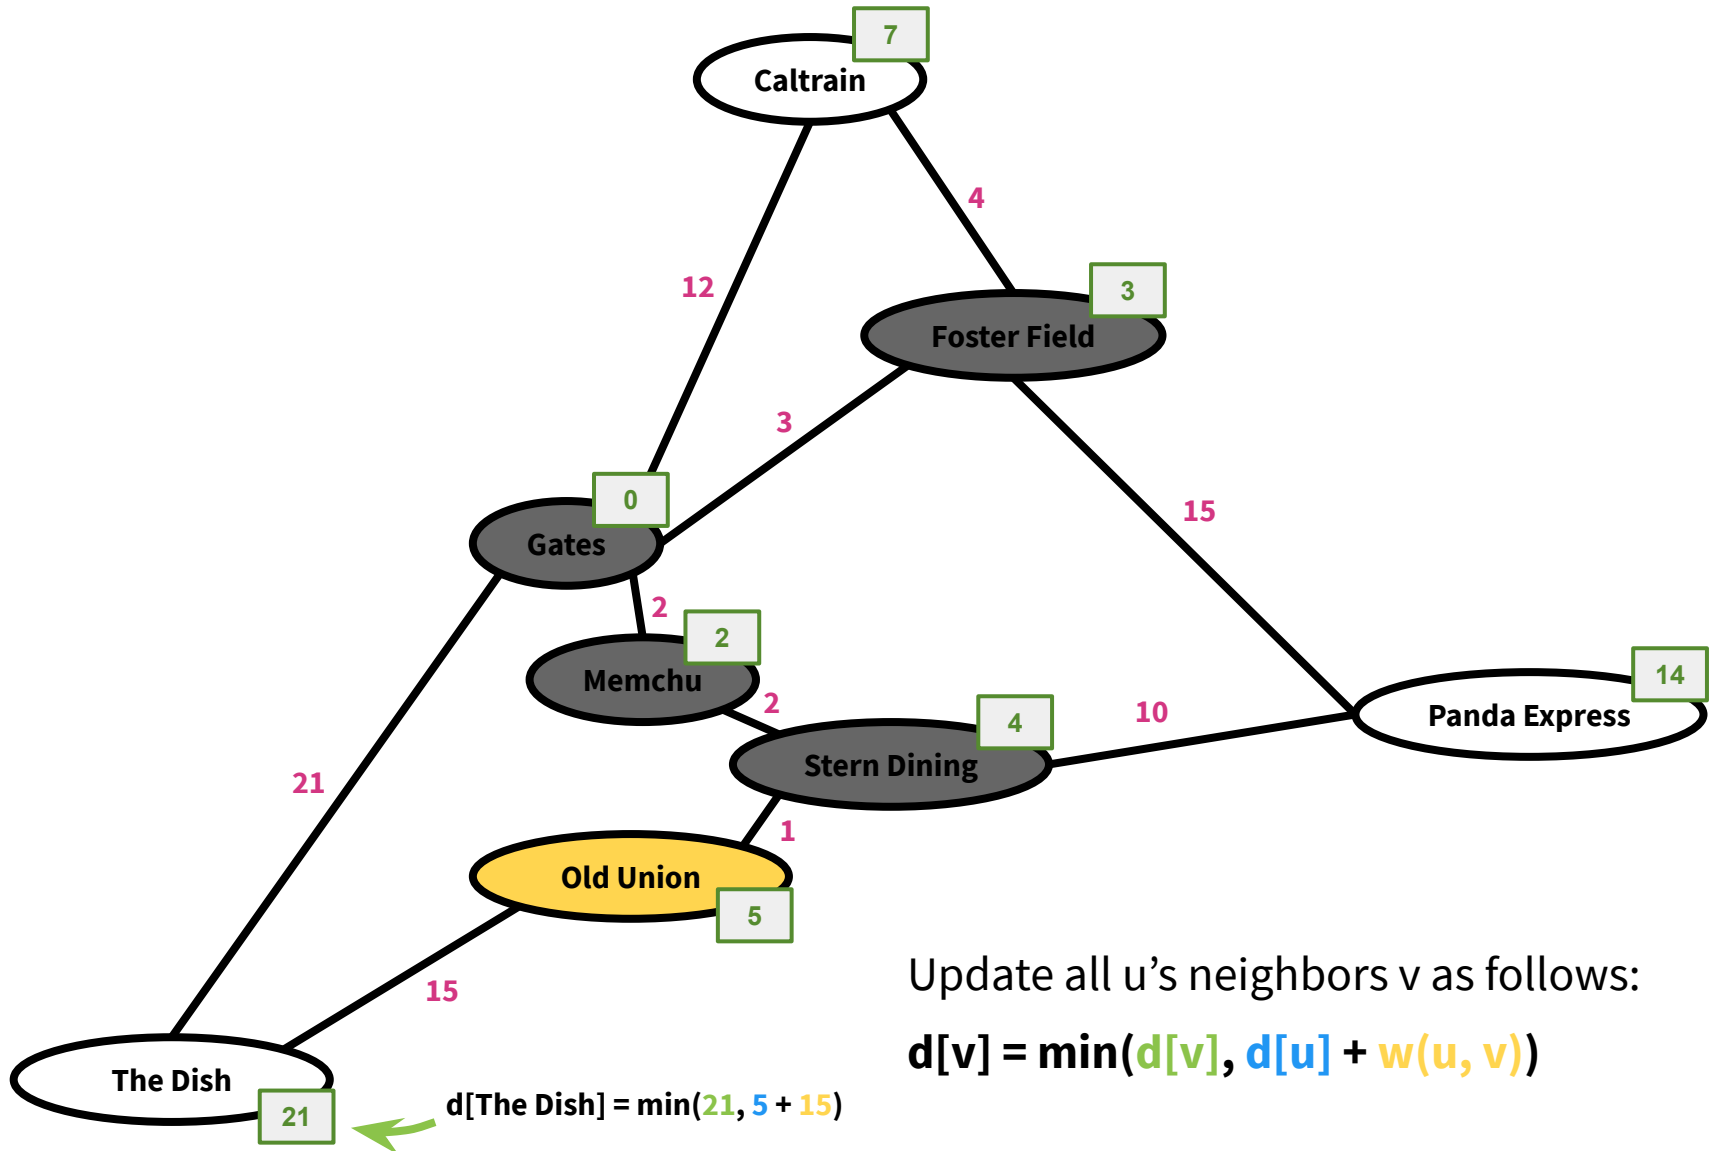
Dijkstra's Algorithm



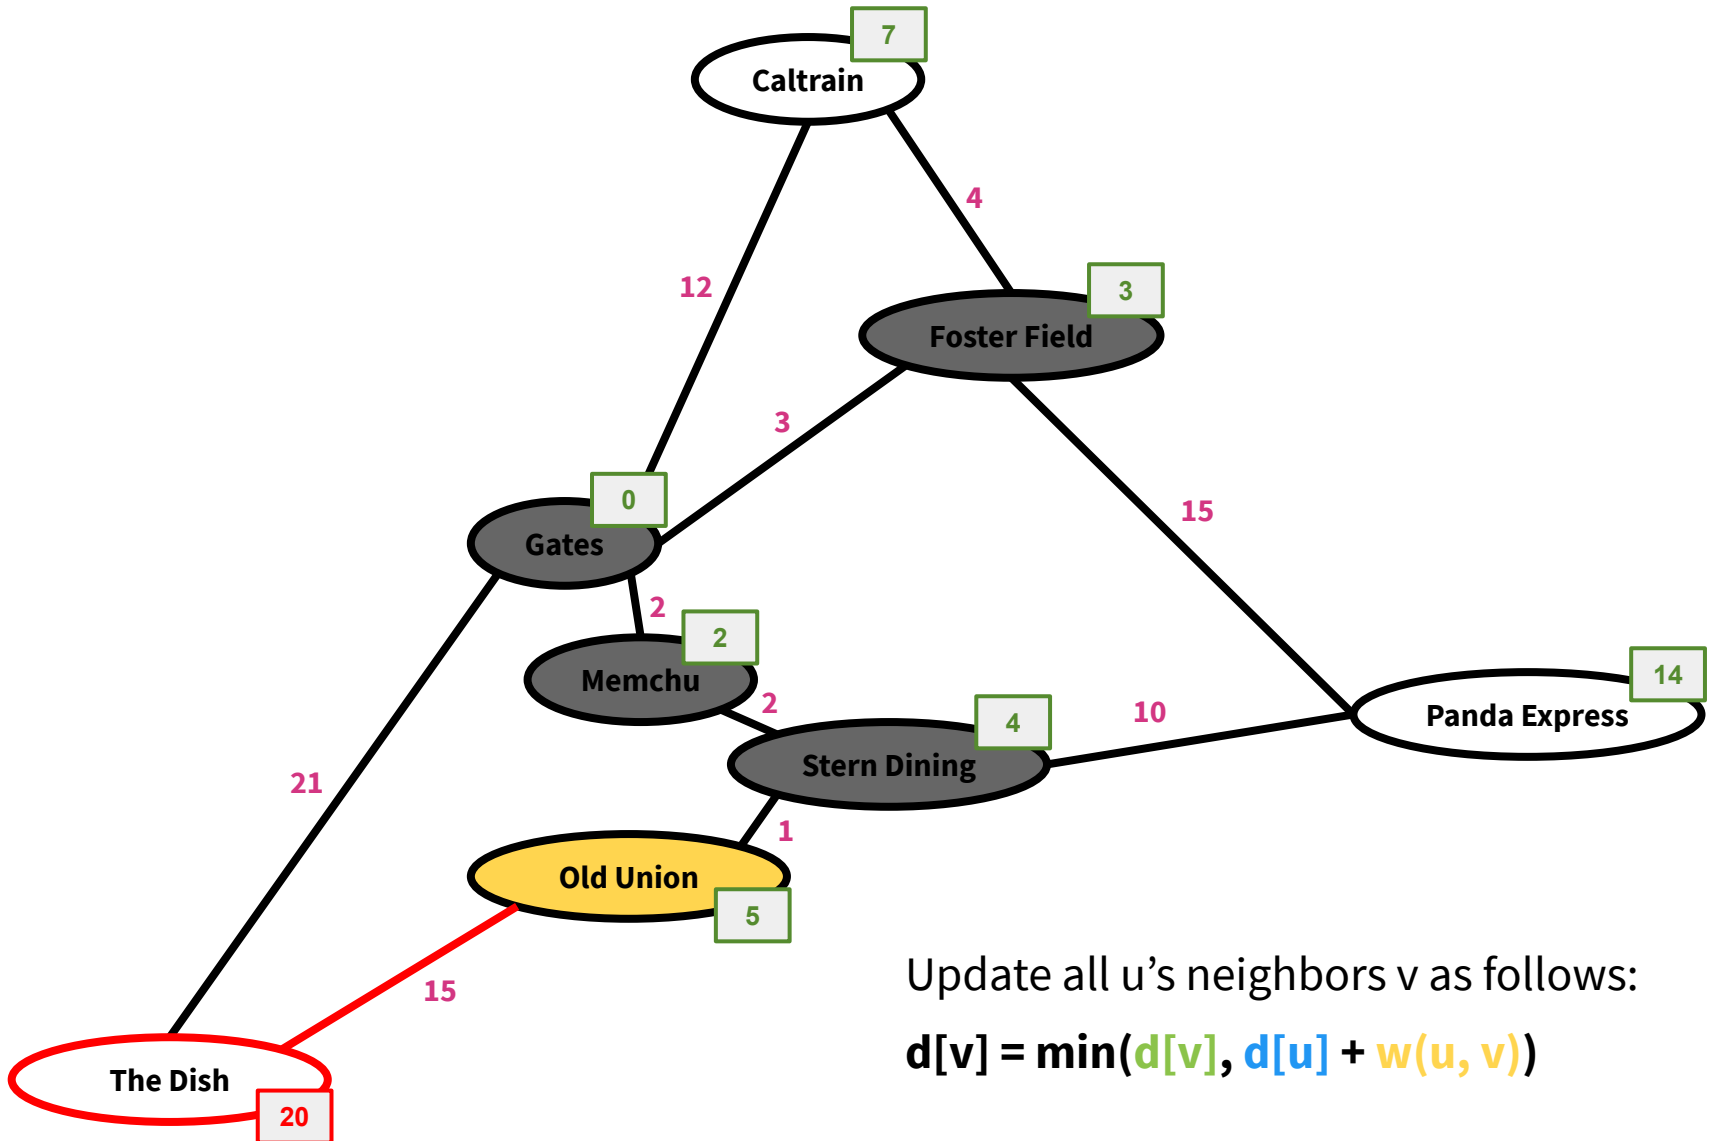
Dijkstra's Algorithm



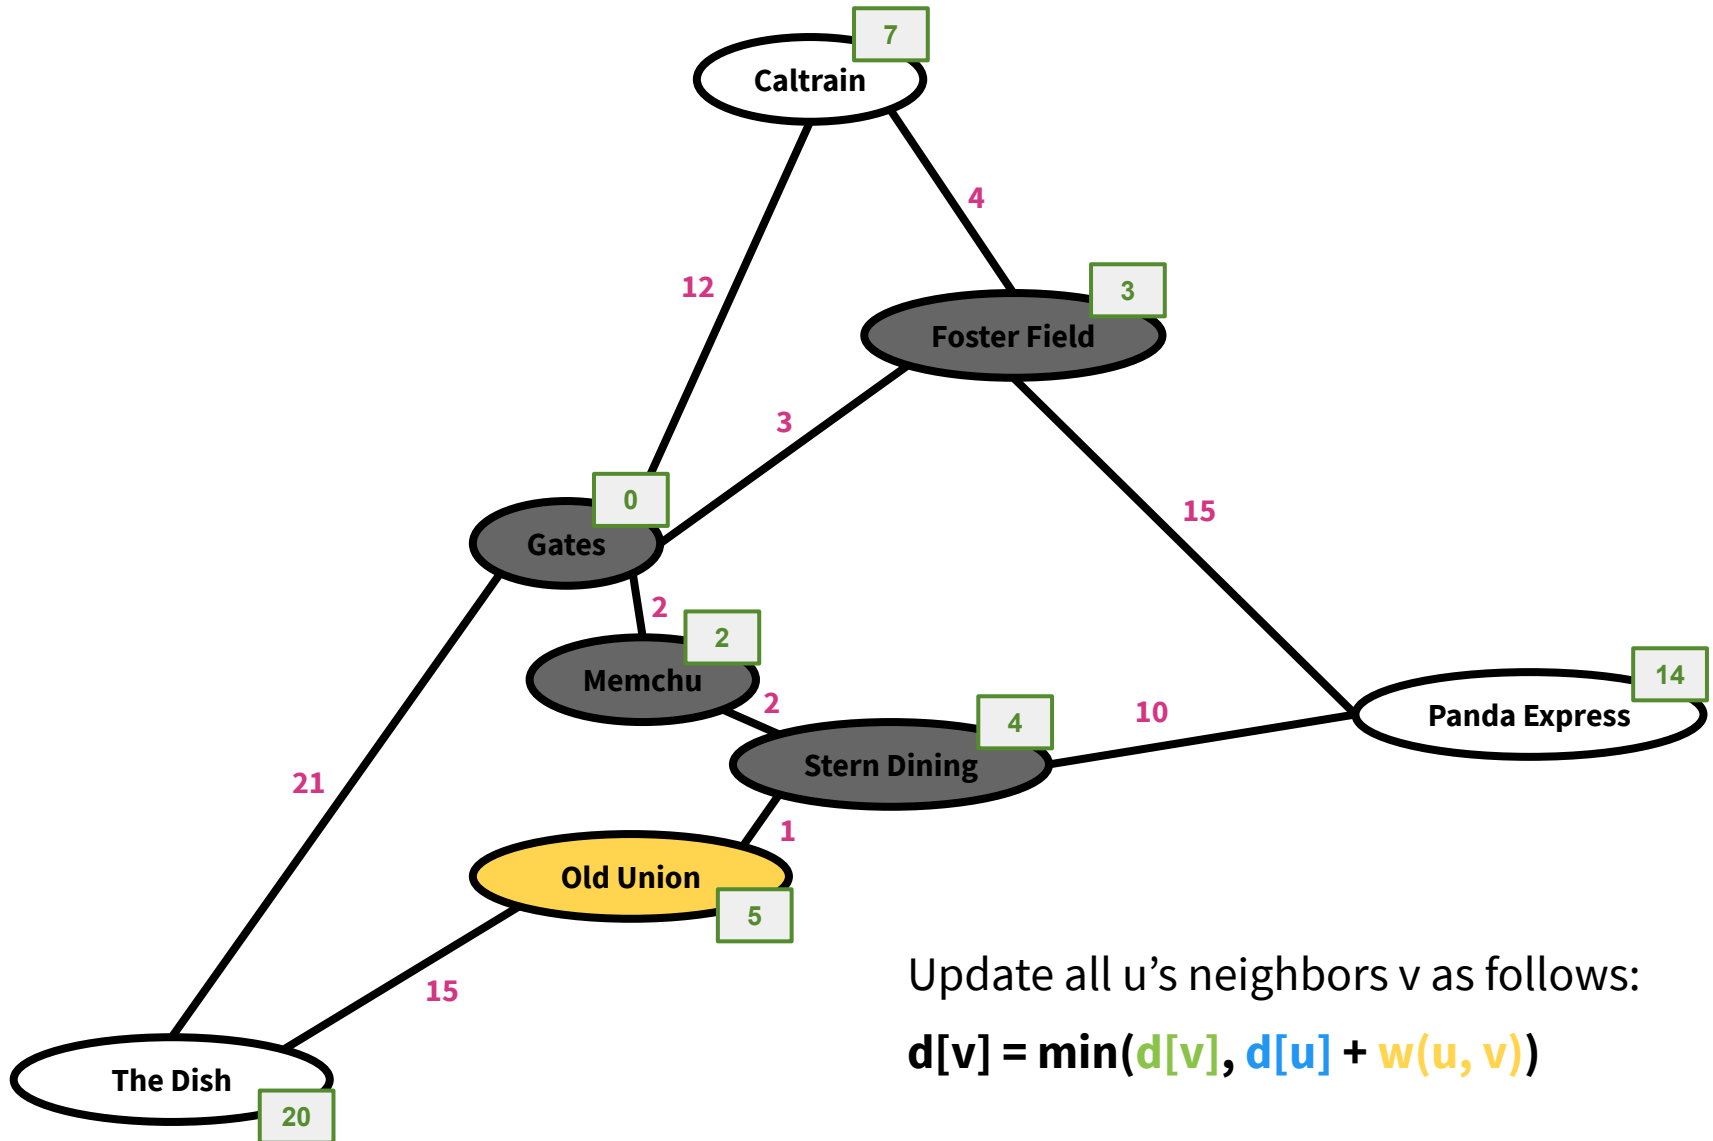
Dijkstra's Algorithm



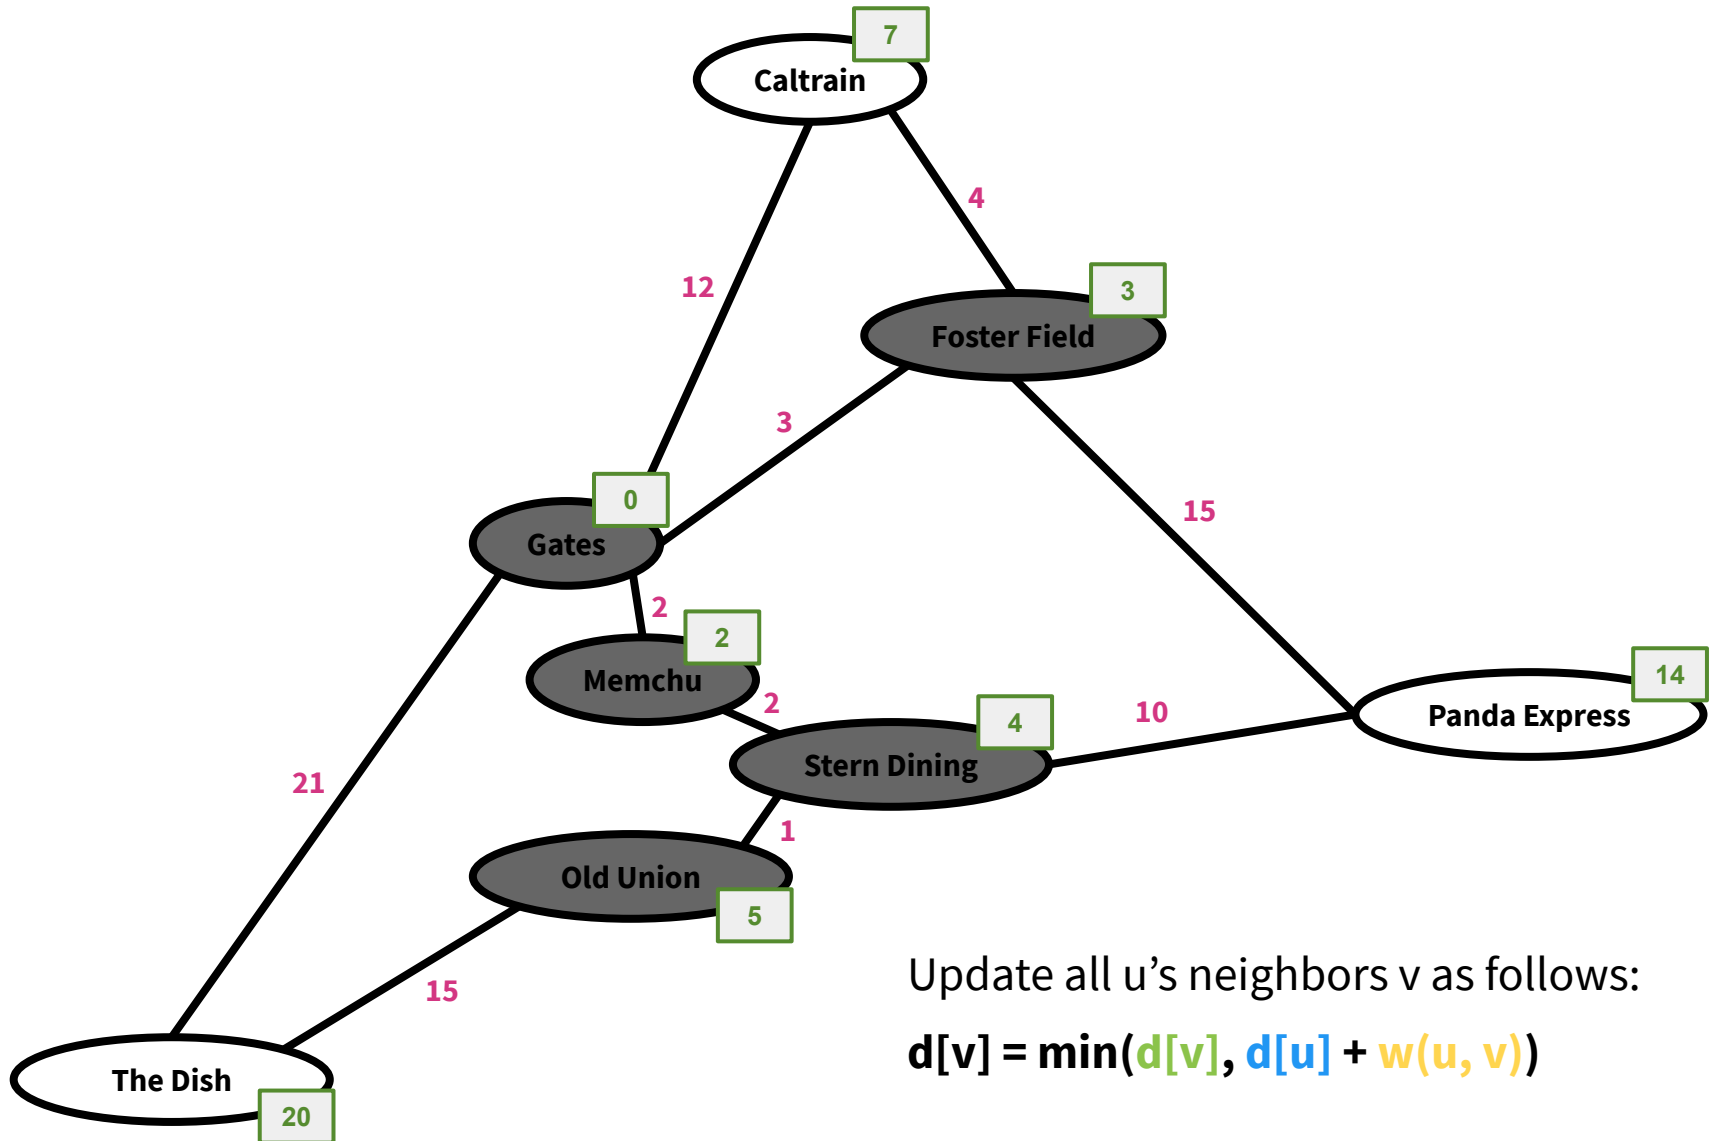
Dijkstra's Algorithm



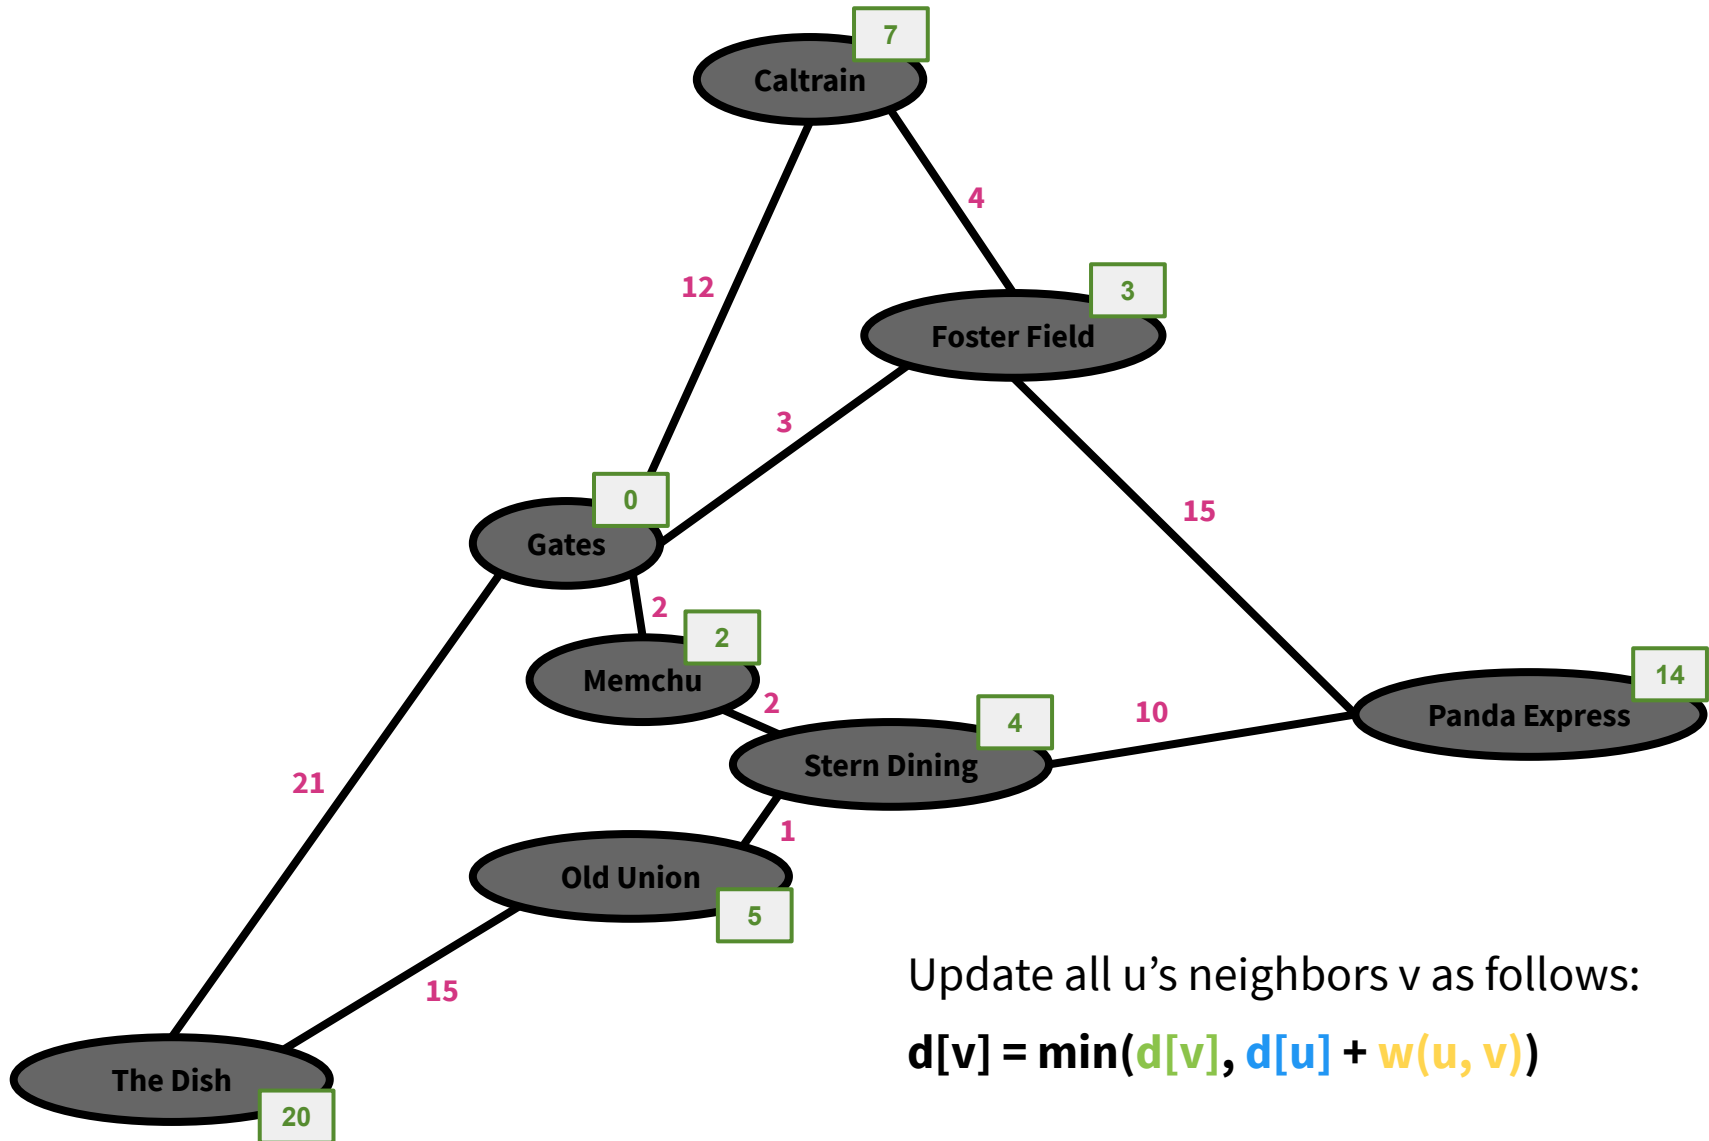
Dijkstra's Algorithm



Dijkstra's Algorithm



Dijkstra's Algorithm



Dijkstra's Algorithm

Why does this work?

Let s be the single source.

Theorem: After running Dijkstra's Algorithm, the estimate $d[v]$ is the actual distance $d(s, v)$.

Proof Outline:

Claim 1: For all v , $d[v] \geq d(s, v)$.

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Together, claims 1 and 2 imply the theorem.

Dijkstra's Algorithm

Why does this work?

Let s be the single source.

Theorem: After running Dijkstra's Algorithm, the estimate $d[v]$ is the actual distance $d(s, v)$.

Proof Outline:

Claim 1: For all v , $d[v] \geq d(s, v)$.

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Together, claims 1 and 2 imply the theorem.

$d[v]$ never increases, so **Claim 1** and **2** imply that $d[v]$ weakly decreases until $d[v] = d(s, v)$ then never changes again.

Dijkstra's Algorithm

Why does this work?

Let s be the single source.

Theorem: After running Dijkstra's Algorithm, the estimate $d[v]$ is the actual distance $d(s, v)$.

Proof Outline:

Claim 1: For all v , $d[v] \geq d(s, v)$.

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Together, claims 1 and 2 imply the theorem.

$d[v]$ never increases, so **Claim 1** and **2** imply that $d[v]$ weakly decreases until $d[v] = d(s, v)$ then never changes again.

By the time we're sure about “done” about v , $d[v] = d(s, v)$.

Dijkstra's Algorithm

Why does this work?

Let s be the single source.

Theorem: After running Dijkstra's Algorithm, the estimate $d[v]$ is the actual distance $d(s, v)$.

Proof Outline:

Claim 1: For all v , $d[v] \geq d(s, v)$.

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Together, claims 1 and 2 imply the theorem.

$d[v]$ never increases, so **Claim 1** and **2** imply that $d[v]$ weakly decreases until $d[v] = d(s, v)$ then never changes again.

By the time we're sure about “done” about v , $d[v] = d(s, v)$.

All vertices are eventually “done” (stopping condition in algorithm).

Dijkstra's Algorithm

Why does this work?

Let s be the single source.

Theorem: After running Dijkstra's Algorithm, the estimate $d[v]$ is the actual distance $d(s, v)$.

Proof Outline:

Claim 1: For all v , $d[v] \geq d(s, v)$.

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Together, claims 1 and 2 imply the theorem.

$d[v]$ never increases, so **Claim 1** and **2** imply that $d[v]$ weakly decreases until $d[v] = d(s, v)$ then never changes again.

By the time we're sure about “done” about v , $d[v] = d(s, v)$.

All vertices are eventually “done” (stopping condition in algorithm).

Therefore, all vertices end up with $d[v] = d(s, v)$.

Dijkstra's Algorithm

Why does this work?

Claim 1: For all v , $d[v] \geq d(s, v)$.

Proof:

We proceed by induction on t , the number of iterations completed by the algorithm.

Dijkstra's Algorithm

Why does this work?

Claim 1: For all v , $d[v] \geq d(s, v)$.

Proof:

We proceed by induction on t , the number of iterations completed by the algorithm.

After $t = 0$ iterations, $d(s, s) = 0$ and $d(s, v) \leq \infty$ which satisfy $d[v] \geq d(s, v)$.

Dijkstra's Algorithm

Why does this work?

Claim 1: For all v , $d[v] \geq d(s, v)$.

Proof:

We proceed by induction on t , the number of iterations completed by the algorithm.

After $t = 0$ iterations, $d(s, s) = 0$ and $d(s, v) \leq \infty$ which satisfy $d[v] \geq d(s, v)$.

For the inductive step, suppose the inductive hypothesis holds for iteration t .

Dijkstra's Algorithm

Why does this work?

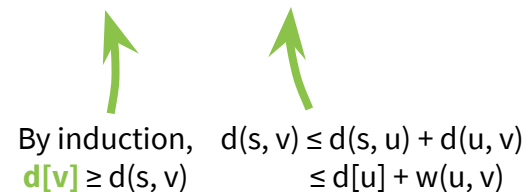
Claim 1: For all v , $d[v] \geq d(s, v)$.

Proof:

We proceed by induction on t , the number of iterations completed by the algorithm.

After $t = 0$ iterations, $d(s, s) = 0$ and $d(s, v) \leq \infty$ which satisfy $d[v] \geq d(s, v)$.

For the inductive step, suppose the inductive hypothesis holds for iteration t . Then at iteration $t + 1$, the algorithm picks a vertex u and for each of its neighbors v sets: $d[v] = \min(d[v], d[u] + w(u, v)) \geq d(s, v)$.



By induction, $d[s, v] \geq d(s, v)$ $d(s, v) \leq d(s, u) + d(u, v)$
 $\leq d[u] + w(u, v)$

Thus, the induction holds for $t + 1$.

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof:

We proceed by induction on t , the number of vertices marked as “done.”

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof:

We proceed by induction on t , the number of vertices marked as “done.”

For the base case, note that after s is marked as “done”, $d[s] = d(s, s) = 0$, which satisfies $d[v] = d(s, v)$.

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof:

We proceed by induction on t , the number of vertices marked as “done.”

For the base case, note that after s is marked as “done”, $d[s] = d(s, s) = 0$, which satisfies $d[v] = d(s, v)$.

For the inductive step, assume that for all vertices v already marked as “done”, $d[v] = d(s, v)$. Let x be the vertex with minimum distance estimate. We must prove $d[x] = d(s, x)$.

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof, cont.:

We proceed by contradiction. Suppose $d[x] \neq d(s, x)$.

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.


Proof, cont.:

We proceed by contradiction. Suppose $d[x] \neq d(s, x)$.



Let p be the shortest path from s to x . There must exist some z on p such that $d[z] = d(s, z)$. Let z be the closest such vertex to x .

We know $d[z] = d(s, z) \leq d(s, x) < d[x]$.

Weights are
non-negative.



Claim 1 implies
 $d(s, x) \leq d[x]$ and
we assumed that
 $d[x] \neq d(s, x)$.



z must exist since,
at the very least, s is
part of the shortest
path, and $d[s] = d(s, s)$.

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.


Proof, cont.:

We proceed by contradiction. Suppose $d[x] \neq d(s, x)$.


Let p be the shortest path from s to x . There must exist some z on p such that $d[z] = d(s, z)$. Let z be the closest such vertex to x .

We know $d[z] = d(s, z) \leq d(s, x) < d[x]$.


Weights are
non-negative.




Claim 1 implies
 $d(s, x) \leq d[x]$ and
we assumed that
 $d[x] \neq d(s, x)$.



z must exist since,
at the very least, s is
part of the shortest
path, and $d[s] = d(s, s)$.



Otherwise, z would be
the vertex with minimum
distance estimate.



Therefore, $d[z] < d[x]$. But this can't be the case. Why not? Since $d[z] < d[x]$ and x is the vertex with minimum distance estimate, z must be already marked “done.”

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof, cont.:

Since z is already marked “done,” the edges out of z , including the edge (z, z') (where z' is also on p) have been relaxed by the algorithm

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof, cont.:

Since z is already marked “done,” the edges out of z , including the edge (z, z') (where z' is also on p) have been relaxed by the algorithm i.e.

$d[z'] \leq d(s, z) + w(z, z') = d(s, z')$ since z is on the shortest path from s to z' and the distance estimate of z' must be correct.

Dijkstra's Algorithm

Why does this work?

Claim 2: When a vertex v gets marked “done”, $d[v] = d(s, v)$.

Proof, cont.:

Since z is already marked “done,” the edges out of z , including the edge (z, z') (where z' is also on p) have been relaxed by the algorithm i.e.

$d[z'] \leq d(s, z) + w(z, z') = d(s, z')$ since z is on the shortest path from s to z' and the distance estimate of z' must be correct.

However, this contradicts z being the closest vertex on p to x satisfying $d[z] = d(s, z)$. Thus, our assumption that $d[z] < d[x]$ must be false, and it follows that $d[x] = d(s, x)$. ■

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .
- There must exist a vertex z closest to x on p for which $d[z] = d(s, z)$. Notice, by our assumption that $z \neq x$.

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .
- There must exist a vertex z closest to x on p for which $d[z] = d(s, z)$. Notice, by our assumption that $z \neq x$.
- But z cannot be the closest vertex to x on p ; simply consider the next vertex z' along the path.

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .
- There must exist a vertex z closest to x on p for which $d[z] = d(s, z)$. Notice, by our assumption that $z \neq x$.
- But z cannot be the closest vertex to x on p ; simply consider the next vertex z' along the path.
 - Since the subpath of a shortest path is also a shortest path, and p is a shortest path from s to z' to x , then the subpath s to z' must also be a shortest path.

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .
- There must exist a vertex z closest to x on p for which $d[z] = d(s, z)$. Notice, by our assumption that $z \neq x$.
- But z cannot be the closest vertex to x on p ; simply consider the next vertex z' along the path.
 - Since the subpath of a shortest path is also a shortest path, and p is a shortest path from s to z' to x , then the subpath s to z' must also be a shortest path.
 - z must have been marked as “done” since it’s on the subpath to x and weights are non-negative.

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .
- There must exist a vertex z closest to x on p for which $d[z] = d(s, z)$. Notice, by our assumption that $z \neq x$.
- But z cannot be the closest vertex to x on p ; simply consider the next vertex z' along the path.
 - Since the subpath of a shortest path is also a shortest path, and p is a shortest path from s to z' to x , then the subpath s to z' must also be a shortest path.
 - z must have been marked as “done” since it’s on the subpath to x and weights are non-negative.
 - Before z was marked as “done”, $d[z']$ was updated to $d[z] + w(z, z')$, which must equal $d(s, z')$ since s to z to z' is a shortest path.

Dijkstra's Algorithm

Another wording of Claim 2

- When a vertex v gets marked “done”, $d[v]$ must be $d(s, v)$.
- By contradiction, assume there exists an x such that when it gets marked as “done,” $d[x] \neq d(s, x)$.
- Consider the shortest path p from s to x .
- There must exist a vertex z closest to x on p for which $d[z] = d(s, z)$. Notice, by our assumption that $z \neq x$.
- But z cannot be the closest vertex to x on p ; simply consider the next vertex z' along the path.
 - Since the subpath of a shortest path is also a shortest path, and p is a shortest path from s to z' to x , then the subpath s to z' must also be a shortest path.
 - z must have been marked as “done” since it’s on the subpath to x and weights are non-negative.
 - Before z was marked as “done”, $d[z']$ was updated to $d[z] + w(z, z')$, which must equal $d(s, z')$ since s to z to z' is a shortest path.
- **Thus, contradiction!**

Bellman-Ford

Bellman-Ford Algorithm

Dijkstra's algorithm solves the single-source shortest path problem in weighted graphs.

Sometimes it works on graphs with negative edge weights, but sometimes it doesn't work.

Bellman-Ford also solves the SSSP problem in weighted graphs.

Always works on graphs with negative edge weights (when a solution exists).


Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

$d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.

We know $k = 0$
i.e. shortest
paths to each
vertex with at
most 0 edges
in it.



	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$				
$d^{(2)}$				
$d^{(3)}$				



Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

$d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.




Bellman-Ford Algorithm

How do we use $d^{(k-1)}$ to fill in $d^{(k)}[b]$?

Recall $d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.


Case 1: the shortest path from s to b with at most k edges actually has at most $k - 1$ edges.



Suppose $k = 3$.

$d^{(k)}[b] = d^{(k-1)}[b]$ i.e. the shortest path of at most $k - 1$ edges is at least as short as any path of at most k edges.

Case 2: the shortest path from s to b with at most k edges really has k edges.



Suppose $k = 3$.


$d^{(k)}[b] = \min_a \{d^{(k-1)}[a] + w(a, b)\}$
i.e. the shortest path of at most k edges is shorter than any path of at most $k - 1$ edges.

Bellman-Ford Algorithm

```
def bellman_ford(G):  
     $d^{(k)}$  = [] for  $k = 0$  to  $|V|-1$   
     $d^{(0)}[v] = \infty$  for all  $v \neq s$   
     $d^{(0)}[s] = 0$   
    for  $k = 1$  to  $|V|-1$ :  
        for  $b$  in  $V$ :  
             $d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\} \}$   
    return  $d^{(|V|-1)}$ 
```

Runtime: $O(|V| |E|)$

Bellman-Ford Algorithm

```
def bellman_ford(G):  
     $d^{(k)}$  = [] for  $k = 0$  to  $|V|-1$   This is a simplification to make the  
     $d^{(0)}[v] = \infty$  for all  $v \neq s$  pseudocode nice. In reality, we'd  
     $d^{(0)}[s] = 0$  only keep two of them at a time.  
    for  $k = 1$  to  $|V|-1$ :  
        for  $b$  in  $V$ :  
             $d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\} \}$   
    return  $d^{(|V|-1)}$ 
```

Runtime: $O(|V| |E|)$

Bellman-Ford Algorithm

```
def bellman_ford(G):  
    d(k) = [] for k = 0 to |V|-1  
    d(0)[v] = ∞ for all v ≠ s  
    d(0)[s] = 0  
    for k = 1 to |V|-1:  
        for b in V:  
            d(k)[b] = min{d(k-1)[b], mina{d(k-1)[a] + w(a,b)} }  
    return d(|V|-1)
```

This is a simplification to make the pseudocode nice. In reality, we'd only keep two of them at a time.

Minimum over all a such that (a, b) ∈ E.

Runtime: $O(|V| |E|)$

Bellman-Ford Algorithm

```
def bellman_ford(G):  
     $d^{(k)}$  = [] for  $k = 0$  to  $|V|-1$   
     $d^{(0)}[v] = \infty$  for all  $v \neq s$   
     $d^{(0)}[s] = 0$   
    for  $k = 1$  to  $|V|-1$ :  
        for  $b$  in  $V$ :  
             $d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\} \}$   
    return  $d^{(|V|-1)}$ 
```



This is a simplification to make the pseudocode nice. In reality, we'd only keep two of them at a time.

Minimum over all a such that $(a, b) \in E$.

Case 1



Runtime: $O(|V| |E|)$

Bellman-Ford Algorithm

```
def bellman_ford(G):  
     $d^{(k)} = []$  for  $k = 0$  to  $|V| - 1$   This is a simplification to make the  
    pseudocode nice. In reality, we'd  
    only keep two of them at a time.  
     $d^{(0)}[v] = \infty$  for all  $v \neq s$   
     $d^{(0)}[s] = 0$   
    for  $k = 1$  to  $|V| - 1$ :  
        for  $b$  in  $V$ :  
             $d^{(k)}[b] = \min\{ \underbrace{d^{(k-1)}[b]}_{\text{Case 1}}, \underbrace{\min_a \{d^{(k-1)}[a] + w(a,b)\}}_{\text{Case 2}} \}$   Minimum over all  $a$   
            such that  $(a, b) \in E$ .  
return  $d^{(|V|-1)}$ 
```

Runtime: $O(|V| |E|)$

Bellman-Ford Algorithm

```
def bellman_ford(G):  
     $d^{(k)} = []$  for  $k = 0$  to  $|V| - 1$   This is a simplification to make the  
    pseudocode nice. In reality, we'd  
    only keep two of them at a time.  
     $d^{(0)}[v] = \infty$  for all  $v \neq s$   
     $d^{(0)}[s] = 0$   
    for  $k = 1$  to  $|V| - 1$ :  
        for  $b$  in  $V$ :  
             $d^{(k)}[b] = \min\{ \underbrace{d^{(k-1)}[b]}_{\text{Case 1}}, \underbrace{\min_a \{d^{(k-1)}[a] + w(a,b)\}}_{\text{Case 2}} \}$   Minimum over all  $a$   
            such that  $(a, b) \in E$ .  
return  $d^{(|V|-1)}$ 
```

Runtime: $O(|V| |E|)$

 Slower than Dijkstra's
 $O(|E| + |V| \log(|V|))$



Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

for $k = 1$ **to** $|V|-1$:

for b **in** V :

$$d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\} \}$$



Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.


for $k = 1$ **to** $|V|-1$:

for b **in** V :

$$d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\}\}$$

We will use table $d^{(0)}$ to fill in $d^{(1)}$. More generally, we will use table $d^{(k-1)}$ to fill in $d^{(k)}$.

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$				
$d^{(3)}$				



Bellman-Ford Algorithm


We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

for $k = 1$ **to** $|V|-1$:

for b **in** V :

$$d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\} \}$$

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$	0	2	4	3
$d^{(3)}$				



Bellman-Ford Algorithm


We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

for $k = 1$ **to** $|V|-1$:

for b **in** V :

$$d^{(k)}[b] = \min\{d^{(k-1)}[b], \min_a\{d^{(k-1)}[a] + w(a,b)\} \}$$

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$	0	2	4	3
$d^{(3)}$	0	2	4	2




Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

Recall $d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$	0	2	4	3
$d^{(3)}$	0	2	4	2




Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

Recall $d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.

The shortest path from s to t with 1 edge has cost ∞ (no path exists).

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$	0	2	4	3
$d^{(3)}$	0	2	4	2



Bellman-Ford Algorithm


We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.

Recall $d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.

The shortest path from s to t with 1 edge has cost ∞ (no path exists).

The shortest path from s to t with 2 edges has cost **3** ($s-v-t$).

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$	0	2	4	3
$d^{(3)}$	0	2	4	2



Bellman-Ford Algorithm

We maintain a list $d^{(k)}$ of length n for each $k = 0, 1, \dots, |V|-1$.


Recall $d^{(k)}[b]$ is the cost of the shortest path from s to b with at most k edges.

The shortest path from s to t with 1 edge has cost ∞ (no path exists).

The shortest path from s to t with 2 edges has cost **3** ($s-v-t$).

The shortest path from s to t with 3 edges has cost **2** ($s-u-v-t$).

	s	u	v	t
$d^{(0)}$	0	∞	∞	∞
$d^{(1)}$	0	2	5	∞
$d^{(2)}$	0	2	4	3
$d^{(3)}$	0	2	4	2



BF Proof of Correctness

We need to prove our main argument.

$d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

Proof: We proceed by induction on k , the number of iterations completed by the algorithm.

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

Proof: We proceed by induction on k , the number of iterations completed by the algorithm.

For our base case, at the start of iteration $k = 1$, the shortest path from s to s with 0 edges has cost 0. The path from s to all vertices $v \neq s$ contains at least 1 edge; there doesn't exist a path from s to v with 0 edges, and this path costs ∞ . Therefore, $d^{(0)}$ is correct.

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

Proof: We proceed by induction on k , the number of iterations completed by the algorithm.

For our base case, at the start of iteration $k = 1$, the shortest path from s to s with 0 edges has cost 0. The path from s to all vertices $v \neq s$ contains at least 1 edge; there doesn't exist a path from s to v with 0 edges, and this path costs ∞ . Therefore, $d^{(0)}$ is correct.

For our inductive step, assume that at the start of iteration k , $d^{(k-1)}[b]$ is the cost of the shortest path from s to b with at most $k - 1$ edges. We consider two cases:

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

Proof: We proceed by induction on k , the number of iterations completed by the algorithm.

For our base case, at the start of iteration $k = 1$, the shortest path from s to s with 0 edges has cost 0. The path from s to all vertices $v \neq s$ contains at least 1 edge; there doesn't exist a path from s to v with 0 edges, and this path costs ∞ . Therefore, $d^{(0)}$ is correct.

For our inductive step, assume that at the start of iteration k , $d^{(k-1)}[b]$ is the cost of the shortest path from s to b with at most $k - 1$ edges. We consider two cases:

Case 1: $d^{(k-1)}[b] < \min_a \{d^{(k-1)}[a] + w(a, b)\}$. This corresponds to the case in which the shortest path contains fewer than k edges. Then our algorithm correctly sets $d^{(k)}[b] = d^{(k-1)}[b]$.

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

Proof: We proceed by induction on k , the number of iterations completed by the algorithm.

For our base case, at the start of iteration $k = 1$, the shortest path from s to s with 0 edges has cost 0. The path from s to all vertices $v \neq s$ contains at least 1 edge; there doesn't exist a path from s to v with 0 edges, and this path costs ∞ . Therefore, $d^{(0)}$ is correct.

For our inductive step, assume that at the start of iteration k , $d^{(k-1)}[b]$ is the cost of the shortest path from s to b with at most $k - 1$ edges. We consider two cases:

Case 1: $d^{(k-1)}[b] < \min_a \{d^{(k-1)}[a] + w(a, b)\}$. This corresponds to the case in which the shortest path contains fewer than k edges. Then our algorithm correctly sets $d^{(k)}[b] = d^{(k-1)}[b]$.

Case 2: $d^{(k-1)}[b] \geq \min_a \{d^{(k-1)}[a] + w(a, b)\}$. This corresponds to the case in which the shortest path contains exactly k edges. Then our algorithm correctly sets $d^{(k)}[b] = \min_a \{d^{(k-1)}[a] + w(a, b)\}$, which minimizes the sum of the shortest path with at most $k-1$ edges to an in-neighbor of b and the weight from a to b .

BF Proof of Correctness

Lemma: $d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges.

Proof: We proceed by induction on k , the number of iterations completed by the algorithm.

For our base case, at the start of iteration $k = 1$, the shortest path from s to s with 0 edges has cost 0. The path from s to all vertices $v \neq s$ contains at least 1 edge; there doesn't exist a path from s to v with 0 edges, and this path costs ∞ . Therefore, $d^{(0)}$ is correct.

For our inductive step, assume that at the start of iteration k , $d^{(k-1)}[b]$ is the cost of the shortest path from s to b with at most $k - 1$ edges. We consider two cases:


Case 1: $d^{(k-1)}[b] < \min_a \{d^{(k-1)}[a] + w(a, b)\}$. This corresponds to the case in which the shortest path contains fewer than k edges. Then our algorithm correctly sets $d^{(k)}[b] = d^{(k-1)}[b]$.

Case 2: $d^{(k-1)}[b] \geq \min_a \{d^{(k-1)}[a] + w(a, b)\}$. This corresponds to the case in which the shortest path contains exactly k edges. Then our algorithm correctly sets $d^{(k)}[b] = \min_a \{d^{(k-1)}[a] + w(a, b)\}$, which minimizes the sum of the shortest path with at most $k-1$ edges to an in-neighbor of b and the weight from a to b .

At the start of iteration $k = |V|$, the algorithm terminates and $d^{(|V|-1)}$ is correct.

BF Proof of Correctness

We need to prove our main argument.

$d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges. 

What else to do? 

BF Proof of Correctness

We need to prove our main argument.


$d^{(|V|-1)}[b]$ is the cost of the shortest path from s to b with at most $|V|-1$ edges. 🙌

What else to do? 🤔

We still need to prove that this argument implies `bellman_ford` is correct
i.e. $d^{(|V|-1)}[a] = \text{distance}(s, a)$.

To show this, we'll prove that the shortest path with at most $|V|-1$ edges is the shortest path with any number of edges (if a shortest path exists).

If the graph has a negative cycle, a shortest path might not exist!




BF Proof of Correctness

But if there's no negative cycle.

There's always a simple shortest path.

A simple path has
no cycles.



BF Proof of Correctness


But if there's no negative cycle.

There's always a simple shortest path.

A simple path has
no cycles.



How do we know this cycle
doesn't help? 🤔




BF Proof of Correctness

But if there's no negative cycle.

There's always a simple shortest path.

A simple path has
no cycles.



How do we know this cycle
doesn't help? 🤔 Since there's no
negative cycles!




BF Proof of Correctness

But if there's no negative cycle.

There's always a simple shortest path.

A simple path has
no cycles.




How do we know this cycle
doesn't help? 🤔 Since there's no
negative cycles! So we remove it.

BF Proof of Correctness

But if there's no negative cycle.


There's always a simple shortest path.

A simple path has
no cycles.



How do we know this cycle
doesn't help? 🤔 Since there's no
negative cycles! So we remove it.

A simple path in a graph with $|V|$ vertices has at most $|V|-1$ edges in it.




BF Proof of Correctness

But if there's no negative cycle.


There's always a simple shortest path.

A simple path has
no cycles.



How do we know this cycle
doesn't help? 🤔 Since there's no
negative cycles! So we remove it.

A simple path in a graph with $|V|$ vertices has at most $|V|-1$ edges in it.



We can't add another edge to this
s-t path without making a cycle
(an edge from s to b wouldn't be
along the path).

BF Proof of Correctness

Theorem: `bellman_ford` is correct as long as the graph has no negative cycles.

Proof:

By our lemma, $d^{(|V|-1)}[b]$ contains the cost of the shortest path from s to b with at most $|V|-1$ edges. If there are no negative cycles, then the shortest path must be simple, and all simple paths have at most $|V|-1$ edges. Therefore, the value the algorithm returns, $d^{(|V|-1)}[b]$, is also the cost of the shortest path from s to b with any number of edges.

Bellman-Ford Algorithm

Bellman-Ford gets used in practice.

e.g. Routing Information Protocol (RIP) uses it. Each router keeps a table of distances to every other router. Periodically, we do a Bellman-Ford update.

Dynamic Programming

Bellman-Ford is an example of **dynamic programming**!

Dynamic programming is an algorithm design paradigm.

Often it's used to solve optimization problems e.g. **shortest** path.